

SPPA: A Runtime Contract and Preliminary Prototype for Semantic Primitive Proxy Actors in UAV Digital Twins

Authors omitted for review

Abstract

Telemetry-driven UAV digital twins can receive object-level perception messages that are too sparse to render as recognizable actors directly. This creates a semantic telemetry-to-actor gap: the perception stream may provide a class, confidence, pose estimate, bounding region, or footprint, while the rendering layer needs an immediately visible object with meaningful occupied volume and predictable cost. This paper specifies and preliminarily instruments Semantic Primitive Proxy Assembly (SPPA), a bounded representation contract for compiling semantic labels into low-polygon proxy actors made of boxes, cylinders, ellipsoids, cones, and tori. SPPA is not a planner occupancy map, not a certified detect-and-avoid layer, not a 3D reconstruction method, and not a replacement for video. It is a proposed actor-representation layer within a larger UAV digital-twin pipeline. The current artifacts include a legacy seven-class OBJ/MTL debug prototype, a bounded-label Python resolver/fallback path, a Python SPPA-DESC-0.2/SPPA-UPD-0.2 descriptor and update-packet implementation, and a native Unreal backend-switch plus descriptor-ingestion smoke test, Editor-Cmd actor microbenchmark, and component payload replay; they do not yet consume live YOLO tracks inside Unreal or provide VR/user validation. We provide the runtime contract, descriptor schema, semantic compiler specification, evidence-status table, preliminary local stress tests against text/image-to-3D generators, lifecycle timing measurements, descriptor/update-packet measurements, a packaged internal Unreal render benchmark, and an evaluation roadmap. The artifacts do not yet establish dense VR frame rate, operator readability, fallback behavior, comparative bandwidth advantage, or end-to-end UAV operation.

1 Motivation

The target semantic-telemetry UAV visualization architecture replaces continuous video transmission with lightweight semantic telemetry. The UAV does not transmit every pixel of the camera stream; it transmits object-level meaning: class label, confidence, position, timestamp, and approximate geometry. The ground station reconstructs the operational scene inside a Cesium-enabled Unreal Engine digital twin for a remote pilot wearing VR goggles.

This architecture creates a practical rendering problem. A detector may report many classes: people, cyclists, cars, trucks, animals, vegetation, poles, barriers, construction machinery, unknown obstacles, and rare objects. A digital twin cannot realistically maintain a complete, manually curated 3D asset library for every object that might appear during real-world UAV operation. Even if a large library is available, asset retrieval, licensing, semantic matching, level-of-detail management, and runtime loading remain operational problems.

Modern image-to-3D and text-to-3D models address a different objective. They are designed to create visually rich 3D assets from images or prompts. This is valuable for content creation, simulation authoring, games, and design. However, remote UAV piloting has a different priority: a detected object must appear in the pilot’s synthetic world under a predictable rendering budget. In the current local debug path, SPPA emits coarse proxy meshes in milliseconds; this is a measured prototype property, not a claim that every neural generator necessarily requires seconds under every configuration.

2 Core Idea

The proposed system performs semantic primitive proxy generation. The key claim is simple:

For telemetry-driven UAV visualization, the digital twin may not need to reconstruct the true mesh of every detected object. The hypothesis is that a low-latency, low-polygon, semantically recognizable volume can preserve enough operational context to justify further operator-facing evaluation.

Instead of producing a detailed mesh, the system assembles predefined geometric primitives into a proxy structure. The proxy is not only selected by class name; it is selected by a lightweight semantic and morphological interpretation of that class. If the detector reports `cow`, the system can infer that the object is an animal, a mammal, and a quadruped. That immediately constrains the proxy to include a body volume, a head volume, four legs, and an orientation hypothesis. If the detector reports `truck`, the system can infer a long vehicle body, a cab, wheels, and an elongated footprint. The silhouette then adjusts the proportions: a long truck and a short truck share the same semantic template, but the cargo body is scaled differently. For example:

- a cow may be represented by ellipsoids for body and head, cylinders for legs, and cones for horns;
- a cyclist may be represented by two torus wheels, a simple bicycle frame, and a human torso/head proxy;
- a tree may be represented by a cylinder trunk and clustered green volumes;
- a car or truck may be represented by boxes and wheel tori;
- a bush may be represented by overlapping green ellipsoids.

The exact geometry is intentionally crude. Its purpose is to encode:

- approximate occupied volume;
- object class recognizability;
- spatial location;
- scale relative to the scene;
- extremely low rendering complexity.

Scale and proportion adaptation are only partially implemented, but the implemented vehicle path is no longer a root-scale deformation of the whole actor. For supported vehicle archetypes, explicit metric dimensions drive a semantic part-layout rule: a longer truck extends the cargo/body span and axle placement while cab and wheel dimensions remain bounded by their own part priors. The Unreal update path likewise rejects shape updates that contain only a new global scale and requires replacement part parameters for `shape_param_update`. The generator now also supports a calibrated mask/bbox route: if an explicit image-to-ground scale is supplied, the oriented footprint can derive metric length/width and drive the same parametric part layout. Without that calibration, masks remain image-space evidence only. Real observed-silhouette part fitting remains unevaluated. Non-vehicle archetypes still use template priors or coarse fallback dimensions. This is important for UAV piloting: when flying over an object, the relevant question is not whether the rendered actor is visually beautiful, but whether its occupied volume and part semantics are approximately correct for navigation and operator interpretation. That claim remains a validation target until the planned volume-error and operator studies are completed.

SPPA: language-assisted part decomposition compiled into deterministic 3D primitives

Language can draft the recipe offline; the runtime path uses a reviewed cache entry and telemetry only.

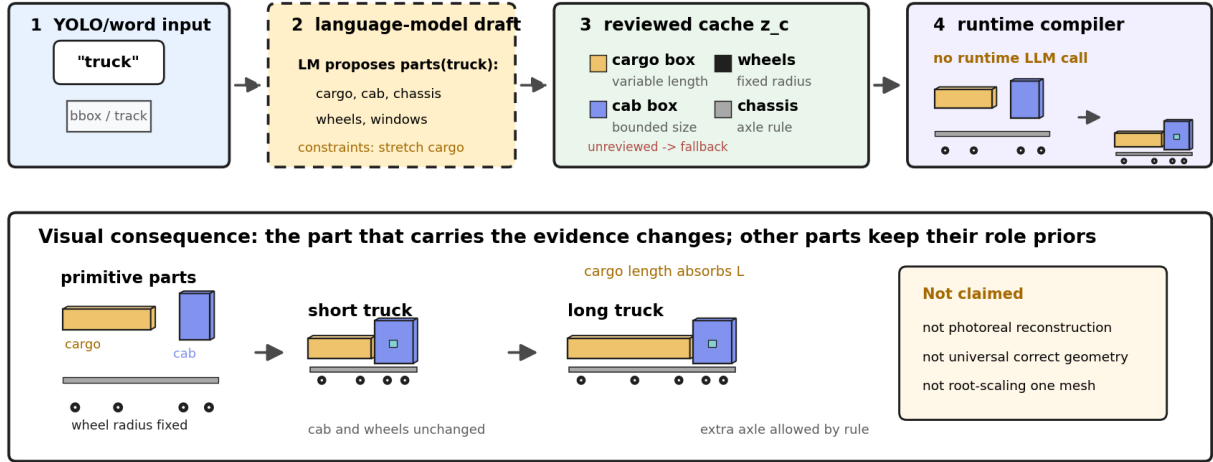


Figure 1: SPPA language-to-parts-to-geometry recipe example. A language model, ontology tool, or human author may draft candidate parts and constraints offline, but runtime uses a reviewed cached descriptor z_c , not a language model or neural mesh generator. The truck example shows the intended chain explicitly: semantic label or word, offline part proposal, reviewed part-role cache, primitive constraints, and deterministic assembly with role-specific scale rules rather than global object scaling.

3 Operational Scope and Non-Claims

The contribution claimed in this manuscript is intentionally narrow. SPPA is a runtime representation contract for intended display proxy actors, not the flight system itself. The surrounding UAV visualization stack provides the broader telemetry, networking, tracking, Unreal/Cesium integration, and VR interface context. SPPA addresses only the question of how an object-level semantic state can be represented visually when a curated asset is unavailable, too costly, or insufficiently matched to the observed object.

The testable claim is therefore:

For telemetry-driven UAV digital twins, track-persistent semantic primitive actors may provide a bounded, uncertainty-aware, intended display obstacle volume when image evidence is missing, degraded, or insufficient for high-fidelity asset generation.

This paper does not claim evaluated UAV/VR teleoperation performance. It also does not claim planning-level risk guarantees, collision avoidance, network-efficiency benefit, or operator benefit. Those are evaluation targets, not results.

SPPA’s novelty is not procedural geometry itself; it is the telemetry-to-runtime-rendering contract with track persistence, uncertainty metadata, and optional fallback beside curated asset spawning.

4 Contributions

This manuscript is framed as a runtime-contract paper with preliminary instrumentation. The contributions are separated by evidence status:

- 1. Implemented template and fallback artifacts.** A seven-class fixed-template debug prototype exports SPPA-style proxy meshes through an OBJ/MTL batch path. This artifact

demonstrates visual style and local generation timing only; it does not implement live telemetry ingestion. Descriptor-level scale evidence and optional native Unreal descriptor ingestion are implemented, including actor-level pose and shape update smoke tests and a packaged synthetic render benchmark up to 100 objects; dense VR/live runtime behavior remains unevaluated. Fallback behavior is smoke-tested in Python/Unreal but not evaluated operationally.

2. **Implemented Python descriptor/update contract.** The current generator emits SPPA-DESC/SPPA-UPD v0.2 artifacts with an explicit semantic resolver contract, bbox/mask/world-pose evidence, yaw provenance and ambiguity, scale source, material source, primitive parts, cost fields, and scheduler metadata. A companion SPPA-UPD-0.2 packet separates create/regenerate payloads from per-frame pose or shape updates. Unreal now includes an optional native descriptor-ingestion path that constructs primitive components from descriptor `parts[]` and an update-packet path for pose/no-op/shape updates while preserving the existing asset backend as the default.
3. **Preliminary measurements.** We report local timing and mesh complexity measurements for the current template path, descriptor/update packet timing and size measurements, a track-lifecycle replay over available obstacle logs, a packaged internal Unreal render benchmark, and stress tests against selected fast text/image-to-3D generators under a constrained PyTorch allocator budget. These are preliminary operating-point measurements, not statistical or end-to-end UAV/VR validation.
4. **Specified runtime contract.** We define the SPPA descriptor, semantic-to-archetype compiler, part-graph representation, uncertainty metadata, yaw ambiguity handling, and uncertainty-marked generic fallback requirements.
5. **Positioning and evaluation protocol.** We position SPPA as a track-persistent semantic actor representation, distinct from primitive shape theory, occupancy mapping, curated asset loading, and neural 3D asset generation. We provide falsifiable hypotheses and a planned evaluation roadmap for the missing Unreal, geometry, operator, fallback, and bandwidth evidence.

5 Research Question and Hypotheses

The scientific question behind SPPA is not whether primitive geometry can look realistic. It is:

What is the minimum semantic-volumetric representation that preserves enough operational information for a future remote UAV visualization interface?

This question leads to falsifiable hypotheses. None of H1–H5 is claimed as evaluated by the current prototype:

H1 latency.

SPPA is designed to reduce detection-to-visible-actor latency relative to runtime asset loading or neural generation. Required baselines are box, capsule, low-poly asset loading, and image-to-3D; primary metrics are P50/P95 latency, triangles, and descriptor bytes. This hypothesis is unevaluated in the current artifact.

H2 recognition.

SPPA is hypothesized to increase class-family recognition relative to generic boxes, capsules, or label billboards at comparable rendering cost. Required baselines are 3D box, capsule/ellipsoid, and billboard label; primary metrics are recognition accuracy, response time, and confusion matrix. A user study is absent.

H3 proportionality.

SPPA is designed to preserve footprint and occupied volume better than fixed-scale templates when object proportions vary. Required baselines are fixed template, bbox-only scale, and curated low-poly asset; primary metrics are BEV IoU, volume error, and height error. This hypothesis is unevaluated in real or simulated UAV detections.

H4 dense runtime.

SPPA is hypothesized to scale to dense VR scenes with more predictable rendering cost than detailed mesh assets. Required baselines are boxes, capsules, low-poly assets, and detailed meshes; primary metrics are FPS, CPU/GPU frame time, and draw calls. A dense-runtime benchmark is absent.

H5 fallback.

SPPA is designed to fall back to uncertainty-marked generic volumes under uncertain class, scale, silhouette, or yaw evidence. Required baselines are no rendering, box, and uncertainty volume; primary metrics are containment rate, false-confidence events, and operator interpretation. A fallback-interpretation study is absent.

6 A Different Runtime Contract from 3D Generation

Recent 3D generative methods have advanced rapidly. DreamFusion introduced a text-to-3D optimization approach using a pretrained 2D diffusion model and score distillation sampling [25]. Magic3D improved quality and speed with a coarse-to-fine strategy, but still reports mesh creation times on the order of tens of minutes [17]. Point-E reduced generation time substantially, producing point clouds in roughly 1–2 minutes on a single GPU, while explicitly trading quality for speed [20]. Shap-E further generates implicit 3D functions from text or image conditions in a matter of seconds [15]. In this manuscript, Point-E and Shap-E are treated as legacy direct text/tag baselines, not as the strongest 2026 image-conditioned asset generators.

More recent feed-forward image-to-3D methods are faster. TripoSR reports single-image mesh reconstruction in under 0.5 seconds [34]. InstantMesh reports diverse single-image 3D assets within approximately 10 seconds [42]. Stable Fast 3D reports textured mesh reconstruction from a single input image in approximately 0.5 seconds [29, 30]. LGM, CRM, and Unique3D are also relevant fast image-conditioned generators [31, 37, 38]. TRELLIS/TRELLIS.2 and Hunyuan3D 2.x focus on high-quality, versatile, textured 3D asset generation using large-scale latent or diffusion architectures [41, 40, 46, 32]. Recent benchmark work, including HY3D-Bench and P3D-Bench, also makes clear that a credible SOTA ranking needs shared reference inputs, explicit metrics, and part/structure evaluation rather than a local qualitative grid assembled from heterogeneous input types [33, 43].

These methods are impressive, but they are not optimized for the same operating point. In UAV VR teleoperation, several constraints become dominant:

- **Per-object latency:** dynamic detections must appear in the digital twin under a fixed runtime budget; neural asset generation may be unsuitable unless it is precomputed, cached, or wrapped in its own tracking/fusion layer.
- **Scene density:** the scene may contain hundreds of people, animals, vehicles, trees, or temporary obstacles.
- **VR frame rate:** the pilot interface must remain stable and comfortable; geometric complexity must be predictable.
- **Input availability:** a clean cropped image suitable for image-to-3D generation may not be available under poor visibility, motion blur, compression, occlusion, or low-confidence detection.

- **Operational sufficiency:** collision avoidance and situational awareness require approximate volume, not photorealistic surface detail.

The proposed primitive-proxy system therefore does not compete with neural asset generation in fidelity. It occupies a different region of the design space: measured millisecond-scale local proxy construction, low-polygon geometry, and class-conditioned volumetric representation.

Table 1: Operating point of SPPA compared with nearby representation choices.

Method family	Primary out-put	Optimized for	Limitation for UAV/VR
Bounding box	cuboid	speed	weak semantic readability
Curated asset	mesh actor	visual quality	incomplete library, fixed scale
Image-to-3D	detailed mesh	fidelity	latency and mesh complexity
3D detection	3D box/state	localization	not an intended operator display actor
SPPA	part proxy	operational volume	coarse geometry

7 Relation to 3D Detection, Primitive Representations, and Human Factors

The proposed approach is related to several existing research directions. The claim is not that primitive part-based representation is new in isolation. The claim to be tested is narrower: SPPA is a proposed rendering/backend layer that turns semantic telemetry into bounded, low-polygon proxy actors for a UAV digital twin when exact assets or dense reconstruction are unavailable.

7.1 Monocular 3D Object Detection

Monocular 3D object detection predicts 3D properties such as object depth, dimensions, rotation, and oriented 3D bounding boxes from a single image. Omni3D introduced a large benchmark with 234k images, more than 3 million instances, and 98 categories, together with Cube R-CNN as a unified detector across diverse camera and scene types [3]. Recent work also studies open-vocabulary monocular 3D detection, aiming to lift 2D detections into metric 3D space beyond closed class vocabularies [44].

These methods are highly relevant because they can provide exactly the parameters that semantic primitive proxies need: approximate depth, orientation, and dimensions. However, their output is normally a 3D bounding cuboid or object state, not a semantically readable, part-based VR actor. Our proposal can use monocular 3D detection as an upstream estimator, then convert the detected class and dimensions into a low-polygon operational proxy.

7.2 Animal Pose and Animal 3D Reconstruction

For animals, the state of the art includes keypoint-based pose estimation and model-based reconstruction. AP-10K provides a benchmark for mammal pose estimation across 23 animal families and 54 species [45]. 3D-Fauna reconstructs articulated 3D meshes of quadruped animals from a single image within seconds using a pan-category animal model [16]. More recent work such as SAM 3D Animal uses prompts such as masks and keypoints to reconstruct multiple animals in the wild using SMAL-based priors [13].

These systems demonstrate that richer animal reconstruction is possible. The operational question is whether such reconstruction is necessary for UAV teleoperation. For a pilot, the

immediate need is often not fur texture, exact limb articulation, or animatable mesh quality. The immediate need is: a large quadruped occupies this volume, its body is oriented in this direction, and this space should be avoided. A primitive proxy can represent that information with substantially lower geometric and runtime cost.

7.3 Primitive and Superquadric Decomposition

Representing objects as combinations of simple shapes is a long-standing idea in computer vision and robotics. Superquadrics are especially relevant because they compactly represent boxes, cylinders, spheres, ellipsoids, and related forms with few parameters. Earlier shape-representation work used generalized cylinders and volumetric components to describe object structure [18], and recognition-by-components proposed that objects can be recognized from arrangements of simple geon-like parts [2]. Superquadric recovery and parsing further developed compact primitive models for object abstraction [28, 24]. Learned volumetric primitive abstraction also shows that primitive assemblies can capture consistent part structure across shape collections [35], while recent systems such as SuperDec explicitly use superquadric primitives as a compact representation for 3D scene decomposition [10].

This literature is a direct prior, not a minor analogy. SPPA must therefore be judged as an operational telemetry-to-rendering contract rather than as a new theory of primitive shape representation. The remaining open question is whether a constrained primitive representation is useful in the specific semantic-telemetry setting, compared with simpler boxes/capsules and with existing asset or mapping representations.

7.4 Procedural Actors, LOD, and Scene Graphs

SPPA is also related to procedural modeling, shape grammars, level-of-detail systems, impostors, semantic scene graphs, and bounding-box visualization. These are closer baselines for SPPA’s runtime contract than generic asset-generation systems. Procedural modeling and CGA shape grammars generate structured geometry from compact rules [19]. ShapeAssembly similarly uses executable programs to assemble cuboid part proxies and represent families of related shapes [14]. Against this literature, SPPA’s proposed difference is not procedural generation itself; it is the use of a bounded procedural descriptor as an online telemetry rendering fallback.

Real-time graphics has long used hierarchical level of detail, impostors, and extreme simplification to control rendering cost [4, 5]. Modern engines and geospatial standards continue this line through virtualized geometry, instancing, and tiled streaming representations [8, 23]. SPPA must therefore be benchmarked against low-poly assets, LOD, impostors, and instanced primitive rendering rather than only against neural 3D asset generators.

Semantic scene graphs are another direct neighbor. 3D scene graphs and dynamic scene graphs connect objects, semantics, geometry, dynamics, and actionability [1, 27]. SceneGraphFusion shows that incremental 3D semantic graph prediction can run online from RGB-D sequences [39]. SPPA can be interpreted as a rendering/backend layer for such semantic state, not as a replacement for scene-graph perception. This framing reduces the novelty claim but makes the system boundary clearer.

7.5 Human Factors, Display Uncertainty, and Mapping Baselines

Because the intended output is for a remote pilot, SPPA cannot rely on visual plausibility alone. Situation awareness and workload should be evaluated with established human-factors instruments such as SAGAT and NASA-TLX [6, 7, 11]. Any claim about operator readability, obstacle-avoidance usefulness, or workload reduction is therefore remains a hypothesis until controlled operator studies are performed.

SPPA also overlaps with robotic mapping representations. Occupancy grids, OctoMap, TSDFs, and ESDFs provide spatial volumes and obstacle-distance information for planning and navigation

[12, 22]. These representations are not intended to be semantic proxy actors by themselves, but they are essential baselines for any claim about operational volume. The publishable evaluation must therefore distinguish operator-facing semantic readability from collision or planning utility.

The closest practical comparison is therefore not “SPPA versus high-fidelity 3D generation,” but SPPA versus no rendering, 3D boxes, capsules/ellipsoids, billboard labels, curated low-poly assets, procedural actors, occupancy-style volumes, and instanced rendering in the target Unreal/VR environment.

Table 2: Nearest alternatives that should be evaluated before claiming an operational advantage for SPPA. The current manuscript specifies this comparison but does not yet execute the Unreal/VR baselines.

Alternative	What it provides	Why it is a required baseline
3D box / capsule / ellipsoid	Minimal occupied volume and very low render cost	Tests whether semantic part graphs add recognition or volume value.
Billboard label	Fast semantic display with almost no geometry	Tests whether class text alone is enough for operator interpretation.
Curated low-poly asset	Hand-authored recognizable actor	Tests the cost and coverage tradeoff against asset libraries.
Unreal ISM/HISM primitive batches	Native instancing of repeated geometry	Tests draw-call and dense-scene scaling in the target engine.
HLOD / impostor / Nanite asset	Engine-level LOD or virtualized geometry	Tests whether standard graphics tooling already solves the runtime-cost problem.
OctoMap / TSDF / ESDF / occupancy volume	Planning-oriented obstacle volume	Tests containment, clearance, and false-confidence tradeoffs.
3D/dynamic scene graph	Semantic object-state representation	Tests whether SPPA is best framed as a rendering backend for graph state.
3D Gaussian splat / scene asset	Visual scene representation	Tests fidelity when a richer 3D scene representation is available.
Text/image-to-3D mesh	Generated asset from prompt or crop	Tests whether neural generation can replace bounded proxy actors under missing or degraded evidence.
SPPA	Track-persistent semantic primitive actor descriptor	Target method; currently specified and partially prototyped.

8 Operation Inside a Semantic-Telemetry UAV Digital Twin

SPPA is intended to run at the ground station or inside the rendering component of the digital twin as a subcomponent of a semantic-telemetry visualization stack, not as the flight system or telemetry architecture itself. The upstream pipeline detects objects, estimates track state, and provides semantic/geometric evidence. SPPA consumes that state only when an entity must be represented, updated, or reparameterized as a display proxy actor.

The runtime loop is:

1. The UAV detects an object and estimates its class, confidence, image region, and geospatial position.
2. The ground station associates the detection with an entity track.
3. SPPA selects or updates a semantic part graph for that entity.
4. The silhouette, footprint, and temporal track adjust orientation, scale, and part proportions.
5. Unreal Engine instantiates or updates primitive components using the resulting proxy descriptor.
6. The interface renders a low-polygon proxy actor for the object in VR; whether this is stable and interpretable for a pilot remains an empirical question.

This loop separates perception from visual fidelity. If a high-quality asset is available and affordable, it may be used. If not, SPPA is intended to provide a proportional, semantically meaningful proxy volume; this remains an evaluation target rather than an operational guarantee.

9 Proposed Runtime Contract

SPPA is specified as a fallback or primary lightweight rendering layer inside a semantic-telemetry UAV digital twin. Its inputs are semantic detections and geometric estimates from the perception pipeline. Its intended outputs are low-polygon proxy actor descriptors for the digital twin.

Table 3: Semantic primitive proxy runtime contract.

Stage	Input	Output
Detection	image frame	class, confidence, silhouette/bounding box
Georeferencing	UAV pose + camera model	world position and scale estimate
Proxy selection	object class	primitive template or generic fallback
Primitive scaling	silhouette/footprint/volume	scaled proxy geometry
Actor spawning	proxy mesh + transform	Unreal digital twin actor
Runtime update	tracked entity state	pose, scale, confidence, uncertainty

The expected telemetry fields are:

- stable entity identifier;
- detected class label;
- confidence score;
- geospatial position;
- estimated footprint, height, or bounding volume;
- silhouette or 2D bounding box when available;
- timestamp and sensor source;
- uncertainty estimate.

10 SPPA Specification

SPPA is specified as a deterministic compiler contract. It should not call a large generative model at runtime, and it should not search a large asset library. Its task is to transform the minimum available semantic and geometric evidence into an operational 3D proxy descriptor. The specified pipeline is:

1. **Normalize the semantic label.** The detector output is mapped to a semantic path such as `cow → animal → mammal → quadruped` or `truck → vehicle → cargo vehicle`.
2. **Select a morphological archetype.** The semantic path selects a small family of part graphs, such as `quadruped`, `biped`, `long vehicle`, `vegetation blob`, `tree-like`, `pole-like`, or `generic obstacle`.
3. **Extract geometric measurements.** The system estimates major and minor footprint axes, aspect ratio, height, ground contact, and yaw candidates from the bounding box, mask, footprint, previous track, and UAV camera geometry.

Table 4: Evidence boundary for the current SPPA artifact.

Component or claim	Evidence in this manuscript	Boundary of interpretation
Fixed proxy templates	Seven OBJ/MTL debug templates are implemented and timed.	Shows local export feasibility only; not a live runtime loop.
Descriptor/update contract	SPPA-DESC-0.2/SPPA-UPD-0.2 are implemented in Python and accepted by an optional Unreal actor path.	Supports the runtime contract; not a VR/headset or live-network validation.
Scale and footprint evidence	Bbox, mask, pose, yaw, and explicit metric dimensions can enter descriptor construction; synthetic mask rows are measured.	Real UAV masks, camera calibration, and objective geometry error are not measured.
Materials and uncertainty	Material-role, evidence-source, yaw-ambiguity, and fallback tags are exported and smoke-tested in Unreal.	Operator interpretation and false-confidence effects are untested.
Packaged Unreal path	Synthetic 10/50/100-object packaged render replay and a short three-stream recorded-payload packaged replay are reported with Tick frame times and estimated component/draw-group counts.	No headset, phase-aligned GPU-profiler, live-network, broad curated-asset, or operator result is reported.
Bandwidth positioning	SPPA JSON packet bytes are measured and compared with modeled box, mesh, and compressed-video payloads in a preliminary link-budget artifact.	This is not a measured radio-link replay; SPPA does not beat compact binary boxes on bytes.

4. **Fit pose and proportions.** The algorithm chooses a yaw angle and assigns length, width, and height to semantic parts while respecting class-level priors and observed silhouette proportions.
5. **Assemble primitives.** Each semantic part is instantiated as a low-resolution primitive: box, cylinder, ellipsoid, cone, torus, capsule, or superquadric-like approximation.
6. **Estimate uncertainty and fallback.** If class, orientation, or scale is ambiguous, the output includes uncertainty metadata or falls back to an uncertainty-marked generic volume.

10.1 Semantic Normalization

SPPA should support many class labels without requiring a manually authored asset for each one. The label is therefore mapped to a compact ontology of morphological archetypes. The ontology does not need to be large; it only needs to capture the body plan relevant for operational volume:

Table 5: Examples of semantic normalization and archetype selection.

Detected label	Semantic path	Archetype
cow, horse, dog	animal / mammal / quadruped	quadruped
person, pedestrian	human / biped	biped
biker, cyclist	human + two-wheel vehicle	rider vehicle
car, van	vehicle / wheeled	compact vehicle
truck, bus	vehicle / long wheeled	long vehicle
tree	vegetation / trunk + canopy	tree-like
bush, shrub	vegetation / low blob	vegetation blob
pole, mast, tower	vertical structure	pole-like
unknown	obstacle	generic volume

10.2 Bounded Semantic-to-Part Compiler

The important architectural step in SPPA is the conversion from a detected semantic label to a constrained geometric part graph. This can be viewed as a semantic compiler. The input vocabulary may be open-ended, but the output is restricted to a finite set of morphological archetypes, semantic roles, and primitive geometries.

Let c be the detected label. A language model, knowledge graph, or manually curated ontology maps c to a semantic descriptor:

$$z_c = \mathcal{L}(c) = (h_c, a_c, q_c, d_c),$$

where h_c is a category hierarchy, a_c is a set of expected attributes, q_c is a set of qualitative morphology cues, and d_c contains coarse dimension priors. For example, a descriptor for **cow** may include **animal**, **mammal**, **quadruped**, **has head**, **has torso**, and **four legs**; a descriptor for **truck** may include **vehicle**, **wheeled**, **cab**, **cargo body**, and **elongated footprint**.

The semantic descriptor selects the closest morphological archetype:

$$A^*(c) = \arg \max_{A \in \mathcal{A}} \text{sim}(\phi(A), z_c) \quad \text{s.t.} \quad \text{risk}(A, z_c) \leq \rho_{\max},$$

where $\mathcal{A}$ is a compact library of archetypes, $\phi(A)$ is the semantic signature of an archetype, and the risk constraint prevents over-specific mappings when semantic evidence is weak. The current artifact does not yet implement a learned or calibrated sim / risk model. Instead, it implements a bounded offline review gate: a candidate recipe must use an allowed primitive vocabulary, known material roles, a reviewed runtime archetype or explicit fallback, no runtime LLM dependency, and a bounded part/triangle budget. This is an operational cache-admission rule, not a calibrated open-vocabulary classifier. The functions sim , risk , descriptor calibration, and threshold $\rho_{\max}$ still need empirical specification before claiming open-vocabulary coverage. If no archetype is sufficiently reliable, the compiler selects an uncertainty-marked generic obstacle.

The selected archetype is expanded into a semantic part graph:

$$G_c = \Gamma(A^*, z_c) = (V_c, E_c),$$

with nodes

$$v_i = (r_i, \omega_i, \pi_i, \alpha_i, \kappa_i).$$

Here r_i is the semantic role of the part, ω_i is the primitive type selected from a finite primitive alphabet $\Omega = \{\text{box, cylinder, ellipsoid, cone, torus, capsule}\}$, π_i is the anchor relative to the parent part, α_i stores allowed aspect-ratio ranges, and κ_i stores importance and uncertainty metadata. Edges E_c encode attachment and relative placement constraints, such as head-attached-to-torso, wheels-attached-to-chassis, or canopy-above-trunk.

Finally, the graph is instantiated as geometry:

$$\mathcal{P}(G_c, \theta) = \bigcup_{v_i \in V_c} T(\pi_i, \theta_i) \omega_i(\alpha_i),$$

where θ_i contains the fitted pose and scale parameters of each part. This equation is the architectural core of SPPA: reviewed semantic labels are not converted directly into arbitrary meshes; they are mapped to a bounded, low-polygon graph of semantic parts drawn from a finite primitive vocabulary.

This gives SPPA bounded semantic coverage only after a label has a reviewed morphological descriptor or a reviewed parent archetype. The quality of the proxy depends on the quality of that descriptor and on the available visual evidence. If no reviewed mapping exists, the specified display policy degrades toward an uncertainty-marked generic volume rather than inventing object-specific geometry.

10.3 Part Graphs

Each archetype is represented by a semantic part graph. Nodes are parts and edges encode relative placement constraints. A node contains:

- semantic role, such as torso, head, leg, cab, cargo body, wheel, trunk, canopy, or shaft;
- primitive type, such as box, ellipsoid, cylinder, cone, or torus;
- relative anchor and transform;
- allowed aspect-ratio and scale range;
- recognition importance;
- polygon budget.

For example, a quadruped graph contains a body ellipsoid, a head ellipsoid, a short neck cylinder, four leg cylinders, and an optional tail. A long-vehicle graph contains a cargo box, cab box, wheels, and optional windshield marker. The same graph can represent many real dimensions because part scales are not fixed.

10.4 Measurement Extraction

The SPFA specification is intended to consume different levels of geometric evidence:

- **Footprint available:** fit a minimum-area oriented rectangle and extract length, width, aspect ratio, and yaw candidates.
- **Silhouette mask available:** use principal axes, asymmetry, protrusions, and lower contact regions to refine orientation and part allocation.
- **Bounding box only:** combine image-space extent with class priors, UAV pose, and ground-plane assumptions to estimate a bounded volume.
- **Temporal track available:** use velocity and previous yaw to reduce front/back ambiguity and jitter.

This is a degraded-sensing design target, not an evaluated field result. A mask can provide additional geometric evidence for proxy fitting, but real UAV mask quality and downstream operator benefit remain unevaluated. When only a class label and a bounding region are available, the intended behavior is a bounded proxy volume for evaluation rather than a detailed reconstruction.

10.5 Uncertainty-Marked Fallback

The SPFA specification requires a usable display volume, even when the class is unknown or the silhouette is too weak for reliable part assignment. The fallback is selected from simple geometric evidence:

- elongated and low footprint: oriented capsule or elongated box;
- small footprint and large height: pole-like cylinder;
- low circular footprint: vegetation blob or generic ellipsoid;
- upright human-scale region: biped proxy or vertical capsule;
- ambiguous region: uncertainty-marked oriented bounding volume.

The fallback is not treated as a successful reconstruction. It is an uncertainty-marked display mechanism intended to preserve operational continuity when semantic detail is missing. The operator sees a volume with uncertainty, and the digital twin does not silently drop an obstacle. Whether this reduces risk or creates false confidence is a validation question.

10.6 Pose and Proportion Fitting

SPPA estimates orientation and dimensions by combining weak cues, but the yaw estimate must distinguish axial and directional evidence. Footprint principal axes and silhouette PCA normally provide an unoriented axis modulo π , while velocity, asymmetric part evidence, or a previous signed track yaw can provide a direction modulo 2π . The implementation therefore first estimates an axial orientation

$$\bar{\psi}_t^{\text{axis}} = \frac{1}{2} \text{atan2} \left(\sum_k w_k \sin 2\psi_k^{\text{axis}}, \sum_k w_k \cos 2\psi_k^{\text{axis}} \right),$$

and then assigns the front/back sign only when directional evidence is available:

$$(\psi_t, \text{yaw_ambiguous}) = \text{orient} \left(\bar{\psi}_t^{\text{axis}}, \{\psi_j^{\text{dir}}, w_j\}_j \right).$$

For vehicles, the cabin or velocity may identify the front. For quadrupeds, a head-like protrusion or motion direction may identify the head. If the evidence is weak, `yaw_ambiguous = true` and the proxy remains a symmetric obstacle volume. The exact confidence thresholds for orient are specified but not calibrated.

Part dimensions are fitted analytically when possible. For a long vehicle:

$$L_{\text{cargo}} = \alpha L, \quad L_{\text{cab}} = (1 - \alpha)L, \quad \alpha \in [0.55, 0.80],$$

where L is the measured footprint length. Thus a short truck and a long truck are represented by the same graph but different cargo-body scale. For a quadruped:

$$L_{\text{body}} = \beta L, \quad L_{\text{head}} = (1 - \beta)L, \quad H_{\text{leg}} = \lambda H,$$

with β and λ constrained by the quadruped prior and by the observed silhouette. In an implementation, the fitting can be performed by a small discrete hypothesis search over 8–32 candidates rather than a heavy optimization loop.

10.7 SPPA Objective

The full SPPA target is specified as a finite hypothesis selection problem, not as an expensive continuous optimizer. Given a detection D_t , an archetype A , measurements such as footprint F_t , optional mask M_t , and previous track state x_{t-1} , SPPA constructs a small candidate set $\mathcal{H}(A, F_t, M_t, x_{t-1})$. Each candidate separates global pose $\xi_t = (p_t, \psi_t)$ from part parameters q_t , including part scales, local anchors, and fallback mode. The selected proxy is:

$$(\xi_t^*, q_t^*) = \arg \min_{(\xi, q) \in \mathcal{H}(A, F_t, M_t, x_{t-1})} w_F (1 - \text{IoU}_{BEV}(P_{\xi, q}, F_t)) + w_S (1 - \text{IoU}_{mask}(\Pi(P_{\xi, q}), M_t)) \\ + w_T d_{SE(2)}(\xi, \xi_{t-1}) + w_A \ell_A(q)$$

$$\text{s.t.} \quad \text{tri}(P_{\xi, q}) \leq N_{\text{max}}, \quad \hat{\tau}(P_{\xi, q}; E, H) \leq \tau_{\text{max}}.$$

Here $P_{\xi, q}$ is the candidate primitive proxy, IoU_{BEV} measures footprint overlap in bird’s-eye view, IoU_{mask} measures optional image-mask agreement, $d_{SE(2)}$ penalizes pose jitter, and ℓ_A penalizes proportions outside the selected archetype’s allowed ranges. If M_t is unavailable, $w_S = 0$; if no reliable footprint exists, SPPA uses class priors and uncertainty-marked generic fallback volumes. The runtime term $\hat{\tau}$ is explicitly hardware- and engine-dependent: E denotes the rendering engine/runtime path and H denotes target hardware. It cannot be proven from triangle count alone and must be measured in the target Unreal/VR runtime. The current artifact fixes descriptor-contract defaults for scheduler thresholds and geometry-objective weights under **SPPA-SCHED-0.2**. These defaults make descriptor/update behavior reproducible, but they are not empirically calibrated acceptance criteria. Candidate generation and real-data threshold calibration remain future work.

10.8 Role of a Language Model

A language model is not required to generate geometry at frame rate. In SPPA, its role is better understood as a semantic compiler or cache builder. When a new label appears, the language model can propose a descriptor z_c , a parent archetype, and a minimal part graph. That proposal is then reviewed against the finite primitive vocabulary, polygon budget, and bounded fallback rules of SPPA. Once reviewed and accepted, the mapping can be cached and reused deterministically. The current review-gate artifact **SPPA-RECIPE-REVIEW-0.1** checks candidate recipes against allowed primitives, known material roles, runtime resolver behavior, part/triangle limits, fallback behavior, and disabled runtime LLM use. In a five-candidate audit, it accepted a **dump truck** heavy-vehicle mapping, allowed an **unlisted obstacle** only as fallback, and rejected candidates for **person poster**, **delivery drone**, and **construction crane** when they required false semantics, unreviewed topology/material roles, or a runtime LLM. Prompt/model provenance and empirical calibration for future new recipes remain unspecified.

This distinction matters for the target UAV telemetry loop. The runtime loop should not depend on a slow or nondeterministic text-to-geometry call. The language model may help extend the reviewed vocabulary; the runtime path still executes a bounded geometric program. For example, if the detector reports **camel**, the compiler may map it to a quadruped with a larger torso, long neck, and optional hump. If it reports **excavator**, it may map it to a vehicle-like body with tracks, cabin, and arm proxy. If the proposal is uncertain, SPPA falls back to the nearest parent archetype or to an uncertainty-marked generic volume.

The coverage claim is therefore conditional and operational, not photorealistic. The intended behavior is bounded semantic proxy coverage: exact labels use reviewed templates, recognizable terms may use reviewed keyword archetypes, and unsupported labels receive a generic fallback. This does not guarantee a useful mesh for every object in the world, and it does not yet establish reliable open-vocabulary behavior. The method specifies a constrained proxy-coverage policy, not universal 3D reconstruction.

10.9 Runtime Output

At runtime, the preferred output is not an OBJ mesh. OBJ export is useful for debugging and paper figures, but a deployed Unreal system should receive a compact proxy descriptor that can be rendered using instancing or procedural mesh components:

```
{
  "descriptor_schema": "SPPA-DESC-0.2",
  "semantic": {"raw_label": "truck", "archetype": "truck"},
  "evidence": {"bbox_px": {...}, "mask_hash": "..."},
  "pose": {
    "position_world": {"x": 10.0, "y": 4.0, "z": 0.0},
    "yaw_source": "track_velocity",
    "yaw_modulo": "2pi",
    "yaw_ambiguous": false
  },
  "scale": {"dims_m": {"length": 5.2, "width": 2.3, "height": 2.7}},
  "runtime_policy": {"cache_key": "track-17", "action": "create"}
}
```

The current Python implementation emits this descriptor and a companion SPPA-UPD-0.2 update packet. The Unreal plugin now exposes an optional `ConfigureProxyFromDescriptorJson` path. The existing `UnrealAssets` backend remains the default and ignores SPPA fields; the `SemanticProxy` backend first attempts to consume an embedded descriptor object, or its string-serialized counterpart, from the same `/api/ui/data` obstacle object. It then falls back to the older class/confidence template path if the descriptor is absent or invalid. The smoke test verified

a real SPPA-DESC-0.2 fixture with 12 descriptor parts producing 12 Unreal mesh components, material/evidence/uncertainty tags, invalid-JSON rejection without mutating the current proxy, descriptor reconfiguration, and pose-update packet acceptance for an existing descriptor proxy. This is an integration smoke test, not a dense runtime benchmark; in the component path, invalid descriptor/update payloads fall back to the older class/confidence template path.

In the July 2026 descriptor benchmark, full JSON descriptors measured approximately 6.3–8.2 kB at P50–P95 in the synthetic cases, while update packets measured 0.9–4.7 kB at P50–P95. In a 20,392-observation stratified replay subset from seven controller logs expanded over four predeclared shape thresholds, descriptor bytes were 6.2–6.5 kB at P50–P95 and update packets were 1.0–1.4 kB at P50–P95. These are preliminary JSON serialization sizes, not evidence for bandwidth reduction against video, telemetry boxes, or mesh transfer.

10.10 Pseudocode

```
function SPPA(detection, track_state, ontology, archetypes):
    label = normalize_label(detection.class_label)
    semantic_path = ontology.resolve(label)
    measurements = extract_measurements(detection)

    archetype = select_archetype(semantic_path, measurements)
    part_graph = archetypes[archetype]

    pose = fit_pose(
        measurements,
        previous = track_state[detection.track_id],
        semantic_path = semantic_path
    )

    dimensions = fit_dimensions(
        part_graph,
        measurements,
        confidence = detection.confidence
    )

    params0 = analytic_initialization(part_graph, dimensions)
    params = small_hypothesis_search(
        part_graph,
        params0,
        measurements,
        pose,
        max_candidates = 32
    )

    parts = instantiate_primitives(part_graph, pose, dimensions, params)
    parts = enforce_polygon_budget(parts, detection.distance_to_pilot)

    uncertainty = estimate_uncertainty(
        detection, measurements, pose, dimensions, semantic_path
    )

    if uncertainty.too_high:
        parts = uncertainty_marked_generic_fallback(measurements, semantic_path)
```

```

proxy = make_proxy_descriptor(label, semantic_path, pose, parts, uncertainty)
track_state.update(detection.track_id, proxy)
return proxy

```

11 Part-Fitting Examples

This section illustrates how the SPPA specification is applied to two common object families. The examples are not special cases; they show how a class prior and observed geometry jointly determine a part graph and its proportions.

The class prior is modeled as a semantic part graph. A class label activates both a category hierarchy and a minimal part structure. For instance:

`cow → animal → mammal → quadruped`

From this hierarchy the system can infer that a proxy should contain at least a body, a head, and four leg supports. Similarly:

`truck → vehicle → {cab, cargo body, wheels}`

This makes the reconstruction compositional. The system is not generating an arbitrary mesh; it is fitting a small semantic graph of geometric parts to the detected silhouette and estimated world scale.

Let a detection be represented as:

$$D_t = (c, p, B, s, \gamma, t)$$

where c is the detected class, p is the estimated world position, B is the image-space bounding box or silhouette, s is an approximate world scale, γ is confidence, and t is timestamp. The proxy generator selects a primitive template:

$$T_c = \{(g_i, r_i, m_i)\}_{i=1}^{n_c}$$

where g_i is a primitive geometry type, r_i is its relative transform within the object, and m_i is its semantic material. More explicitly, each primitive belongs to a semantic part:

$$T_c = \{(a_i, g_i, r_i, \theta_i, m_i)\}_{i=1}^{n_c}$$

where a_i is a part label such as torso, head, leg, cab, cargo body, or wheel, and θ_i contains part-specific scale limits and deformation rules. The final proxy separates global pose, local part transform, and primitive dimensions:

$$P_t = \bigcup_{i=1}^{n_c} T_{\text{world}}(p_t, \psi_t) T_i(a_i, \delta_i) \omega_i(q_i),$$

where T_{world} places the object in the digital twin, T_i places part i relative to its parent, ω_i is the selected primitive, and q_i contains the fitted primitive dimensions.

The silhouette is used to estimate a coarse yaw angle and to allocate scale to the most likely parts. For a quadruped, the largest elongated region is treated as the torso candidate; smaller protrusions near one end suggest head position; thin lower protrusions suggest legs. For a truck, the elongated body region sets the cargo body length, while a shorter high-volume region can become the cab. In ambiguous cases, the proxy should preserve uncertainty rather than claiming a precise pose.

This proportional fitting is one of the central differences between a fixed asset lookup and a semantic primitive proxy. A conventional asset library may contain a single truck mesh or several manually authored variants. A neural generator may produce a plausible truck mesh, but it is not

necessarily conditioned on the real detected footprint in a way that enforces a low-polygon, VR-stable obstacle volume. The primitive proxy instead treats the detected silhouette as a constraint on part proportions. The same template can therefore represent both a short utility truck and a long cargo truck by stretching the cargo body while preserving cab and wheel semantics.

This formulation makes four design choices explicit:

1. the proxy is conditioned by semantic class and category hierarchy;
2. the proxy is decomposed into a small graph of semantic parts;
3. the proxy is oriented and scaled by the detected silhouette or estimated footprint;
4. the geometry is intentionally constrained to a small number of primitives.

If the class is unknown, the system can fall back to a generic uncertainty volume, such as a translucent bounding box, ellipsoid, or vertical cylinder. This preserves an explicit display placeholder, but whether it changes operator interpretation or reduces risk is a risk and human-factors question that remains unevaluated.

12 Implementation Status

To avoid overstating the current prototype, Table 6 separates implemented functionality from the full SPPA architecture.

Table 6: Current prototype versus full SPPA architecture.

Component	Current prototype	Full SPPA target
Semantic input	Single text label entered via batch script	Detection message with label, confidence, track ID, pose, bbox, mask, footprint, and uncertainty
Class handling	Exact templates plus keyword archetypes and generic fallback; out-of-template labels no longer abort, but unknown labels receive generic geometry	Reviewed ontology/cache with richer archetype descriptors and uncertainty-marked generic fallback
Geometry scaling	Fixed class-level proportions	Footprint-, silhouette-, and track-conditioned part scaling
Orientation	Template default orientation	Fused yaw from footprint, silhouette, velocity, and previous track
Runtime output	OBJ/MTL files for visual debugging	Compact primitive descriptor for instanced Unreal rendering
Evaluation	Per-object generation timing through batch/OBJ path	Latency, polygon budget, dense-scene FPS, footprint/volume error, recognition, and pilot-task metrics

13 Prototype

We implemented a minimal procedural prototype named **XYT-xabi-yolo-telemetry**. The prototype receives a word from a Windows batch file and exports an OBJ/MTL mesh using only procedural geometry. The current supported classes include:

- cow;
- biker;
- tree;
- car;
- truck;
- bush;
- tractor.

This prototype is intentionally simple. The OBJ/MTL debug export path demonstrates the speed and geometry style of the approach, not a full UAV perception integration. In that debug path, each class is mapped to a fixed lightweight primitive template. The newer **SPPA-DESC-0.2** path accepts offline/CLI evidence such as bbox, mask polygon, world pose, yaw/heading, track id, source image size, and explicit metric dimensions. The Unreal plugin now includes an optional descriptor-ingestion API that builds primitive components from descriptor `parts[]` while leaving the curated-asset backend as the default. Open-label handling is implemented through exact templates, keyword archetypes, and a generic fallback, but that fallback has not been evaluated with operators or dense runtime scenes. In the full semantic-telemetry target, this descriptor path still needs live detections, calibrated geospatial footprints, packaged runtime profiling, and native dense-scene benchmarking.

14 Prototype Timing and Geometry Complexity

The prototype was measured by invoking the actual batch interface used by the operator. The observed wall-clock generation times and exported geometry sizes were:

Table 7: Prototype OBJ generation time and geometry complexity. These values measure the observed debug path only, not a deployed Unreal instancing runtime and not an end-to-end UAV/VR path.

Class	Time (s)	Vertices	Faces	OBJ bytes
biker	0.061	488	396	21,263
bush	0.081	256	210	10,888
car	0.062	448	408	20,268
cow	0.065	612	502	26,533
tractor	0.074	500	450	22,629
tree	0.062	288	236	12,308
truck	0.076	856	786	39,182

These values include Python process startup and file export overhead. They are reported only as debug-path observations. They do not establish the latency of a descriptor-based Unreal implementation, dense-scene rendering, headset frame rate, or detection-to-visible-actor latency. The relevant prototype result is modest: seven fixed low-polygon templates can be exported through the current batch/OBJ path in under 0.1 s per object.

15 Preliminary Fast 3D Generator Stress Test

Because recent image-to-3D systems are now fast enough to be plausible competitors, we ran a preliminary stress test against the subset of fast open-source generators that could be executed

reproducibly in the current Windows workstation stack under a constrained portable-GPU profile. The purpose of this test is not to claim a visual advantage over neural generators or to rank the current SOTA. A real SOTA ranking would need common detection crops or reference images, target meshes or human-preference scores, and a method set that includes contemporary systems such as TRELLIS.2, Hunyuan3D 2.1, SF3D, and TripoSR. SPPA would be a runtime semantic-proxy baseline in that study, not an image-to-3D asset generator. The local test asks only whether runnable generators can replace a bounded low-polygon runtime proxy inside an Unreal-based UAV digital twin.

The benchmark used the same six object classes as the current prototype: `cow`, `biker`, `tree`, `car`, `truck`, and `tractor`. SPPA received the class label. Legacy direct text baselines Point-E and Shap-E received the same prompt/tag field. TripoSR and Hunyuan3D received RGBA proxy crops generated from the same object set. Those crops are local benchmark inputs, not real detection ground truth; a submission-quality ranking should replace them with recorded detection crops/masks or a public benchmark reference set. All neural runs used the same desktop RTX 5090 workstation but with a 6 GB PyTorch CUDA allocator budget. This is not a physical GPU partition; it constrains PyTorch’s caching allocator to a fraction of visible device memory [26]. The current input-provenance manifest explicitly records all six first-row images as synthetic proxy crops and the verifier reports zero ground-truth items. This is a guardrail against treating the visual grid as a detector benchmark before recorded detections or reference annotations exist.

The 6 GB budget was chosen as a portable-flight stress profile rather than an arbitrary constraint. NVIDIA lists 24 GB, 16 GB, 12 GB, and 8 GB GDDR7 memory configurations across the 2026 RTX 50-series laptop family [21]. Epic’s Unreal Engine hardware guidance recommends 8 GB or more graphics RAM for UE5 development [9]. Steam’s May 2026 hardware survey reports VRAM shares of 8 GB (25.89%), 12 GB (12.77%), 16 GB (24.05%), 24 GB (5.22%), and 32 GB (1.20%) [36]. A 32 GB desktop RTX 5090 is therefore not a representative deployment assumption for a laptop that may simultaneously run Unreal, video, telemetry, and the generator.

Table 8: Preliminary July 2026 stress test against fast open-source 3D generators under a 6 GB PyTorch allocator budget. This is a local operating-point stress test over six object classes, not a statistical benchmark or end-to-end UAV/VR validation.

Method	Condition	n	Reps	Input source	Median wall	Triangle range	Peak torch reserved
SPPA	current label only	6	1/object	class label	0.0016 s	488–1,668	0 MB
template							
Legacy text, $K = 16$	Shap-E label only	6	1/object	prompt/tag	2.0080 s	17,884–131,668	6,120 MB
Legacy text + SDF32	Point-E label only	6	1/object	prompt/tag	6.4266 s	1,026–5,454	3,644 MB
TripoSR 128 ³ extraction	warm, clean proxy crop	6	1/object	RGBA image proxy	0.4604 s	19,284–31,012	2,028 MB
Hunyuan3D-2mini Turbo, 5 steps	clean proxy crop	6	1/object	RGBA image proxy	1.5042 s	332,020–1,698,032	6,092 MB max

The ranking supported by these local runs is a runtime task-fit ranking rather than a visual SOTA ranking. Table 9 reports the explicit pass/fail criteria used to separate the SPPA contract claim from image-to-3D visual reconstruction quality.

The full logs, commands, process snapshots, and generated meshes are stored in the benchmark runs directory. The image/proxy run is 20260701_195624; the text/tag run is 20260701_text3d_prompt_baselines. Point-E and Shap-E are reported as legacy direct text/tag baselines. Point-E used the `base40M-textvec` model, 4,096 points, and SDF meshing at grid size 32; the reported time includes mesh conversion because Unreal cannot consume a point cloud as a conventional mesh actor without an additional representation layer. Shap-E used `text300M`, batch size 1, fp16 sampling, and $K = 16$ Karras steps. This is a speed-oriented setting rather than Shap-E’s higher-step notebook example, so it should not be read as best-quality Shap-E.

TripoSR [34] was run at marching-cubes resolution 128 and used a local CPU PyMCubes

Table 9: SPPA runtime task-fit ranking derived from the local July 2026 generator stress runs. This ranking is for the UAV/VR semantic-proxy contract, not for image-to-3D visual SOTA. The six score criteria are: median generation below 33 ms, maximum mesh below 5k triangles, maximum torch reserved memory below 1 GB, no GPU inference, deterministic descriptor contract, and track-updateable actor semantics.

Method	Input	Score	Wall	Max tris	VRAM	Failed criteria
SPPA finite archetype	semantic label/tag	6/6	0.0016s	1,668	0MB	-
TripoSR warm	clean RGBA proxy crop	0/6	0.4604s	31,012	2028MB	lat., tris, mem, GPU, contract, update
Hunyuan3D-2mini Turbo	clean RGBA proxy crop	0/6	1.5042s	1,698,032	6092MB	lat., tris, mem, GPU, contract, update
Shap-E text K=16	prompt/tag	0/6	2.0080s	131,668	6120MB	lat., tris, mem, GPU, contract, update
Point-E text + SDF32	prompt/tag	0/6	6.4266s	5,454	3644MB	lat., tris, mem, GPU, contract, update

fallback because the upstream `torchmcubes` CUDA extension did not compile on the Windows, PyTorch, and CUDA stack used here. Hunyuan3D-2mini Turbo used five inference steps, octree resolution 380, and FlashVDM. The inspected Hunyuan3D local text route first invokes a text-to-image model and then shape generation; if benchmarked, it should be reported as text-to-image-to-3D, not as a direct tag-to-mesh baseline. Stable Fast 3D compiled only under a Python 3.10/Torch 2.8 environment in this workstation and its warm benchmark did not emit a model-load event within 20 minutes, so it is treated as unresolved rather than reported as a failed quality result. SPAR3D installed in the same Python 3.10/Torch 2.8 stack after dependency pinning, but model loading failed with a gated-repository access error for `stabilityai/stable-point-aware-3d`; it is therefore also unresolved in this pass.

For qualitative inspection, both runs also store orthographic front, side, top, and isometric views for each generated mesh in their `views` directories. These images are not used as quantitative evidence. They are an audit artifact for checking whether each method preserves a useful footprint, frontal/lateral extent, and top-down occupancy under the kinds of changing viewpoints encountered during flight. Dense neural meshes are sampled for visualization only; the triangle counts in Table 8 are measured from the original mesh files.

These results narrow the claim. SPPA is not a general high-fidelity 3D asset generator. TripoSR and Hunyuan3D are far more appropriate when visual reconstruction from a usable crop is the target, and Shap-E shows that text-conditioned meshes can be produced in seconds under a fast sampling configuration. However, the preliminary data are consistent with the operational distinction motivating SPPA: class-conditioned primitive proxies can be generated in milliseconds, without GPU inference, and with bounded low-polygon geometry. The open question is therefore not whether SPPA should replace neural generation visually; it should not. The open question is whether bounded primitive descriptors are the right runtime representation when the UAV system needs a persistent track object whose pose, scale, yaw, and uncertainty can be updated every frame under a shared Unreal/VR GPU budget. That question remains open until measured with dense rendered scenes and operator-facing tasks.

16 Bounded Label Resolver And Negative-Control Timings

Following the zero-trust audit, we added a bounded label-resolution path to the prototype and to the Unreal SPPA actor. The implementation now separates three cases: exact implemented labels, keyword-resolved archetypes, and generic unknown-label fallback. This is not a claim that every out-of-template word receives a correct object-specific mesh. It is only a measured claim that out-of-template labels no longer abort the SPPA path and instead produce either an archetype proxy or an uncertainty-marked generic proxy.

A smoke batch with four labels produced the following outcomes: `car` resolved as `exact_class`

with 864 triangles; `ambulance` resolved as `keyword_archetype/light_vehicle` with 864 triangles; `antenna` resolved as `keyword_archetype/vertical_structure` with 90 triangles; and `mystery_object` resolved as `fallback_unknown_label` with 94 triangles. The measured artifact is stored under `experiments/sppa_open_label_smoke`.

We also measured deliberately cheap procedural negative controls. These rows do not define the scientific competition for SPPA, and they should not be read as a visual substitutability test. Their purpose is narrower: they bound the minimum local construction cost of rendering something from sparse telemetry on a flight laptop. Table 10 reports 50 in-memory construction repetitions per class over six labels. OBJ export was performed once per method/class only to check geometry generation.

Table 10: Lightweight negative controls over six classes. These are local Python construction timings, not Unreal frame times and not a visual substitutability claim. Lower dimension error means closer AABB match to the target synthetic dimensions; it does not measure semantic readability or operator utility.

Method	n	P50 build ms	P95 build ms	Triangles	Mean dim. error
Billboard	6	0.0007	0.0011	2–2	0.3333
Box	6	0.0046	0.0058	12–12	0.0000
Capsule proxy	6	0.0585	0.0685	212–212	0.0081
Ellipsoid	6	0.0372	0.0447	144–144	0.0000
SPPA fixed	6	0.2924	0.3563	488–1668	0.2852
SPPA global-scaled	6	0.4575	0.4963	488–1668	0.0000
SPPA parametric	6	0.2902	0.3285	488–1692	0.1589

The result weakens any simplistic speed claim. SPPA is slower than a box or billboard, and zero-AABB-error primitives are trivial when the target is only an outer volume. The defensible argument from this Python debug benchmark is restricted to retaining class-level part structure at low local construction cost. This is not Unreal, packaged, or VR frame-time evidence. Whether the additional part structure helps operator recognition or occupied-volume judgement is untested. For that reason, the supplement does not use a visual comparison against these controls as evidence; visual substitutability is addressed through curated-asset and text/image-to-3D comparisons instead.

We separately measured synthetic within-class scale variants, including compact and long vehicles. This tests only explicit-dimension adaptation, not silhouette fitting from real UAV images. Table 11 separates a trivial global-scaling SPPA baseline from the implemented parametric vehicle rule. Global scaling removes AABB error by definition, but it also scales cab, wheels, windows, and body together. The parametric rule does not optimize for zero AABB error; it preserves bounded part roles while changing the parts that should absorb length changes. A box remains the lower-cost geometry-fit baseline; SPPA must justify its extra triangles through recognizability and part semantics, not by claiming a raw AABB fit win.

Table 11: Synthetic scale-variant benchmark. This measures explicit-dimension adaptation only, not mask fitting or real image adaptation. The vehicle-only rows are the fair test for the implemented parametric vehicle grammar; other archetypes remain template-prior or fallback paths.

Method	n variants	Median dim. error	Median build ms	Triangles
SPPA fixed, all archetypes	12	0.3350	0.2941	488–1668
SPPA global-scaled, all archetypes	12	0.0000	0.4619	488–1668
SPPA parametric, all archetypes	12	0.2170	0.2928	488–1692
Box, all archetypes	12	0.0000	0.0045	12–12
SPPA fixed, vehicles only	4	0.3926	0.4201	864–1668
SPPA global-scaled, vehicles only	4	0.0000	0.6401	864–1668
SPPA parametric, vehicles only	4	0.0223	0.3448	864–1692
Box, vehicles only	4	0.0000	0.0045	12–12

As a controlled part-invariance check, two truck descriptors with equal width/height and different length keep cab scale fixed at [1.656, 2.070, 1.952] m and wheel scale fixed at [0.392, 0.094] m,

while the cargo/body length changes from 2.263 m to 5.188 m. This is still a deterministic layout prior, not evidence that SPPA recovers true truck morphology from imagery.

A separate calibrated-mask benchmark exercises the path from supplied mask footprint to geometry. The benchmark uses 64 synthetic rectangle masks, four lengths, two widths, four axial rotations, two explicit ground-sample distances, and 50 build repetitions per case. Because the ground scale is supplied, the derived length/width target is known by construction; SPPA recovered the metric length/width with 0.000/0.000 m P50/P95 summed error, local geometry build time of 442/480 μ s at P50/P95, descriptor-build time of 297/321 μ s at P50/P95, and descriptor size of 12.9/13.6 kB at P50/P95. The no-calibration control stayed on `template_prior`; masks without metric scale are not silently treated as metric geometry. This benchmark is reported as a contract regression, not as real UAV-mask validation. The standalone truck role-preservation image is retained only as a supporting artifact. It is not included as a result figure here because Figure 1 already carries the visual decomposition argument, while the benchmark values above carry the reproducible evidence.

17 Unreal Backend Switch Smoke Test

The Unreal integration was rebuilt with UE 5.7 and tested with the backend verification script. The smoke test checks four properties: the backend enum contains asset spawning and semantic proxy modes; the default mode is asset spawning; explicit backend selection reaches both modes; and the UI toggle path returns to the opposite mode. The report records the enum values as `UNREAL_ASSETS` and `SEMANTIC_PROXY`.

The same test also spawns the semantic proxy actor for `bike`, `cow`, `tower`, `car`, `ambulance`, `tree`, `antenna`, `mystery_object`, and `unknown`. The generated mesh-component counts were 6, 7, 4, 7, 7, 5, 4, 3, and 3 respectively. The script path is `tools/verify_sppa_backend.ps1`.

This test is important for scope control: SPPA is optional and does not replace the existing asset-spawning path. It is also not a dense-runtime benchmark. The artifacts are `pipeline/logs/sppa_backend_verify_latest.json` and `pipeline/logs/sppa_backend_verify_latest.log`. Native frame-time, draw-call, and headset results are not reported by this smoke test.

18 Recorded Payload Adapter and Unreal Replay

We added a recorded-log adapter that converts sampled `obstacle_ingest` events into the same `/api/ui/data-style obstacles[]` payload shape used by the asset and SPPA backends. The common stream contains entity id, class/type, confidence, pose/geographic fields when present, and source metadata. The SPPA stream preserves those common fields and attaches either `sppa_descriptor` for create/regenerate events or `sppa_update_packet` for update events. A changed-only SPPA stream also suppresses no-op packets.

The command was:

```
python tools\sppa_sota_benchmark\run_recorded_telemetry_payload_replay.py
```

The artifact directory is recorded in the run manifest and in the project evidence note for this replay. The two default log files contained only seven recorded obstacle observations across three payload events and four tracks, with labels `biker`, `cow`, and `tower`. The scheduler produced four create actions, one pose update, and two no-ops. Common asset-compatible payloads totaled 2,762 bytes across the three events. SPPA emit-all payloads totaled 48,188 bytes, and the changed-only SPPA stream totaled 43,132 bytes. Local adapter timing over the seven observations was 257.300/374.500 μ s P50/P95 for descriptor build, 3.600/8.000 μ s P50/P95 for scheduler decision, and 26.200/52.500 μ s P50/P95 for update-packet construction. The current generator separates the per-event `input_hash` from the `descriptor_id`: the latter is now a topology-cache identifier, so pose/no-op updates can reuse cached parts instead of appearing as new geometry.

We then consumed the generated JSONL streams through the same Unreal obstacle-batch ingestion entry point used by both asset and proxy backends in Editor-Cmd:

```
powershell -NoProfile -ExecutionPolicy Bypass -File \
tools\replay_recorded_sppa_payloads.ps1 -SkipBuild \
-Backends "unreal_assets,semantic_proxy,semantic_proxy_instanced"
```

The run accepted all 27 payload applications: three streams, three backends, and three payload events. For the SPPA emit-all stream, final actors/components were 4/4 for `unreal_assets`, 4/44 for `semantic_proxy`, and 1/20 for `semantic_proxy_instanced`. The instanced route reports one managed actor because the live backend uses a single HISM batch actor. P50/P95 apply times for SPPA emit-all were 0.539/1.614 ms for assets, 1.344/3.805 ms for the actor proxy, and 2.394/4.049 ms for the instanced proxy. The common asset-compatible stream also reached all three backends; its final actors/components were 4/4, 4/23, and 1/6 respectively.

This replay supports a limited implementation claim: recorded semantic obstacle rows can be exported as common asset-compatible payloads, SPPA can be attached as an optional extension, and the current Unreal component can ingest those payloads while switching between the asset, actor-proxy, and instanced-proxy backends. By itself, this Editor-Cmd replay does not show live network behavior, packaged runtime, render-thread or GPU frame time, VR performance, detector accuracy, georeferencing accuracy, or operator interpretation. The changed-only stream is accepted in this short replay because Unreal retains prior entity state and the test does not cross the stale-entity timeout; it is not evidence for a long-sequence transport policy.

We then replayed the same JSONL streams through the packaged executable benchmark route with rendered frames, using `tools/benchmark_sppa_packaged_render.ps1` and the recorded-payload file/name command-line inputs documented in the artifact manifest.

The packaged run `20260703T024800Z_packaged_recorded_payload_replay` accepted all 27 payload applications. Frame P95 was 16.551 ms for assets, 6.531 ms for the actor proxy, and 4.809 ms for HISM. Apply P95 was 2.566, 5.670, and 3.246 ms for the same backends. The packaged runner counts active instanced draw groups for HISM estimates; in this short replay the HISM route had median active groups/draws of 6 and median estimated triangles of 7,552. This is rendered packaged-executable evidence for the backend switch and payload contract, but still not live network timing, phase-aligned GPU-profiler timing, headset execution, or operator validation.

19 Evidence-Calibrated Material Cues

Material-cue extension. We implemented a bounded material-cue extension after the reviewer-style risk analysis. It is not texture or PBR asset generation. Each proxy part receives a material role, an evidence-source tag, and an uncertainty-style tag. Python exports these cues through MTL comments and a JSON material manifest. Unreal reflects the same contract through component tags for role, evidence source, and uncertainty style. The cues are semantic priors or explicit unknown fallbacks; they are not claimed to be observed textures unless a future sensor/operator source explicitly provides that evidence.

The local debug-path ablation is retained only as a noncentral engineering check. Across six classes and 50 repeated OBJ/MTL runs per class, flat, class-color, part-material, and low-confidence material-metadata variants had similar Python geometry-build medians, while writing the material manifest added approximately 0.23–0.27 ms at P50–P95 in the measured debug path. This is not a comparative result against alternative 3D representations and does not establish perceptual benefit. The visual material-contact sheet has therefore been removed from the manuscript body; it remains a supporting audit artifact, not a paper figure.

The Unreal smoke test was extended to fail if generated proxy parts lack material-role, evidence-source, or uncertainty-style tags. The latest report records the evidence-material switch, material-policy actor tags, semantic-prior tags for known/archetype classes, and fallback-unknown plus

warning-marker tags for unknown classes. The supporting artifacts are the material-cycle note, the ablation run directory, and the latest backend verification report.

20 Operational Claim To Be Tested

The scientific contribution is not a new high-fidelity 3D generator. The testable claim is instead:

For telemetry-driven UAV digital twins, class-conditioned primitive assemblies may provide a bounded, uncertainty-aware, intended display obstacle volume when a curated asset is unavailable or too costly to render.

This claim has not been evaluated as an operator-facing system. It must be tested through generation latency, polygon and draw-call budget, dense-scene VR frame rate, operator recognition of class and occupied volume, navigation usefulness under degraded video, unknown-class fallback behavior, and bandwidth comparison against encoded video, telemetry-only boxes, and full-mesh transfer.

20.1 Temporal Update Policy

SPPA should not be evaluated as if a complete mesh must be regenerated at every video frame. The intended deployable policy is track-persistent. When a stable detection track appears, the intended runtime policy instantiates or compiles one proxy. During subsequent frames, the runtime should update pose, scale, velocity, uncertainty, and visibility state from the shared telemetry/detection stream. Geometry should be regenerated or swapped only when the class, archetype, dimensions, or confidence state crosses a predefined threshold. For static objects, most frames are expected to be transform updates rather than geometry-generation events. For rigid moving objects, geometry is expected to remain resident while pose and yaw are updated. For articulated or shape-changing objects, the descriptor needs parametric parts or explicit shape-update events. The current Unreal actor rejects shape-update packets that contain only a global scale and requires replacement part parameters for shape updates, avoiding silent root-scale deformation of wheels, cabs, or other semantic subparts. The dense replay calibration and visual consequences of the fixed scheduler thresholds are not reported here.

This policy is part of the comparison with image-to-3D generators. A changing UAV viewpoint does not, by itself, require SPPA to regenerate the object if the object identity and dimensions are stable. A neural image-to-3D baseline can also be generated once and instanced if a stable tracking/fusion layer is available. If it is used directly from changing per-frame crops, repeated generation may be required; if it is generated once, the benchmark must include the tracking/fusion layer, asset cache, decimation/LOD, and uncertainty handling. The fair benchmark must therefore report both one-time generation/instantiation cost and per-frame update cost.

Table 12: Timing taxonomy for the current evidence. The rows measure different artifacts and must not be interpreted as native Unreal/VR frame-time evidence.

Measurement	What it measures	n	Reported result	Not evidence for
SPPA stress-test row	Current batch template path over six labels	6	median 0.0016 s	Unreal spawn/update cost
In-memory template creation	Python procedural object construction	60,000	P50 177.2 μ s	Mesh export or rendering
Debug OBJ/MTL export	File-based prototype export path	600	P50 1.076 ms	Instanced primitive runtime
Policy/state update	Offline track-state update during log replay	436,652	P50 0.4 μ s	Render-thread or GPU cost

20.2 Implemented Descriptor and Update-Packet Benchmark

After the first reviewer-driven correction cycle, the Python generator emits a versioned geometry/track descriptor separate from the material manifest. The descriptor records source label, archetype resolution, bbox/mask/world-pose evidence, yaw provenance, yaw ambiguity, scale source, primitive parts, cost fields, and scheduler thresholds. A companion update packet distinguishes creation, pose updates, shape-parameter updates, topology regeneration, low-confidence drops, and no-ops. This is implemented in Python and benchmarked offline. Native Unreal descriptor ingestion is now implemented as an optional actor path and verified by reflection/smoke tests, but the benchmark in this subsection is still a Python contract benchmark rather than Unreal frame-time.

The smoke set generated descriptors for six exact classes, three keyword-based out-of-template labels, and three unsupported labels. All exact and keyword labels resolved to known archetypes; all unsupported labels produced explicit unknown fallback descriptors. This demonstrates bounded label routing, not universal word-to-3D generation.

Table 13: Preliminary SPPA-DESC-0.2/SPPA-UPD-0.2 benchmark from the large replay artifact 20260702_observed_material_v04_large. Times are Python contract and scheduler measurements, not Unreal frame-time.

Measurement	n	P50	P95	P99
Synthetic descriptor build with parts	160	249.8 μ s	407.7 μ s	463.6 μ s
Synthetic scheduler decision	160	3.7 μ s	5.3 μ s	8.3 μ s
Synthetic update-packet build	160	10.2 μ s	62.6 μ s	65.3 μ s
Synthetic create total, Python only	20	567.8 μ s	1,016.3 μ s	1,134.6 μ s
Synthetic pose/no-op update, no mesh	140	13.9 μ s	63.8 μ s	66.6 μ s
Synthetic full descriptor size	160	9,034 B	11,401 B	11,402 B
Synthetic update packet size	160	1,211 B	8,398 B	8,424 B
Replay descriptor build with parts	81,568	391.8 μ s	515.0 μ s	671.2 μ s
Replay scheduler decision	81,568	9.2 μ s	14.2 μ s	19.2 μ s
Replay update-packet build	81,568	16.8 μ s	75.3 μ s	91.0 μ s
Replay pose/no-op update, no mesh	80,156	26.4 μ s	73.0 μ s	96.7 μ s
Replay update packet size	81,568	1,303 B	6,089 B	6,376 B

The same replay contained one descriptor-build outlier at 360.050 ms, so the distribution should be interpreted through P50–P99 rather than as a hard real-time guarantee. Scheduler sensitivity was evaluated at shape thresholds 0.05, 0.10, 0.20, and 0.30. In the synthetic sequences, shape updates decreased from five events at 0.05 to one event at 0.20–0.30. In the stratified replay subset, seven logs contributed up to 3,000 observations each, with the smallest log contributing 2,392 observations. The resulting 20,392 base observations produced 81,568 threshold-expanded rows. Shape updates decreased from 2,106 events at threshold 0.05 to 738 events at threshold 0.30; creates remained fixed at 353 per threshold and no topology regenerations occurred in this subset. One create-total timing outlier reached 243 ms, so P95/P99 are more informative than the maximum for this Python microbenchmark. This sensitivity must be predeclared for future experiments because threshold choice directly changes update frequency. The replay remains policy-derived from logs, not native Unreal actor instrumentation.

20.3 Preliminary Link-Budget Model

We added a byte-accounting model over the same replay to test whether SPPA can honestly support a bandwidth claim. The artifact is `experiments/sppa_bandwidth/20260703_link_budget_model`. Only SPPA JSON packet bytes are measured from existing descriptor/update artifacts. Box JSON, box binary, mesh transfer, and compressed-video rows are explicit models. This is not a UAV radio-link measurement and not a video encoder benchmark.

Table 14: Preliminary link-budget model over 20,392 replay observations. Values are bytes per second. SPPA rows use measured JSON packet sizes; box, mesh, and video rows are models.

Threshold	SPPA all	SPPA changed	Box JSON	Box binary	Mesh create+xfm	2 Mbps video
0.05	94.0 kB/s	37.4 kB/s	10.7 kB/s	3.27 kB/s	36.3 kB/s	250 kB/s
0.10	87.9 kB/s	32.2 kB/s	10.7 kB/s	3.27 kB/s	36.3 kB/s	250 kB/s
0.20	81.8 kB/s	27.3 kB/s	10.7 kB/s	3.27 kB/s	36.4 kB/s	250 kB/s
0.30	78.5 kB/s	25.8 kB/s	10.7 kB/s	3.27 kB/s	36.5 kB/s	250 kB/s

The result weakens any broad bandwidth claim. SPPA does not beat a favorable binary box telemetry model on bytes, and it is similar to the modeled create-plus-transform mesh baseline under the changed-only policy. It is, however, substantially below modeled continuous compressed video rates (250/625/1250 kB/s for 2/5/10 Mbps). The defensible claim is therefore representation positioning, not bandwidth superiority: SPPA occupies a middle ground between compact semantic boxes and video/mesh transfer while preserving a part-based actor representation. A final communication claim still requires a real network replay or controlled link emulator.

20.4 Unreal Descriptor-Ingestion Smoke Test

After implementing `ConfigureProxyFromDescriptorJson`, the Unreal telemetry plugin was rebuilt with Unreal Engine 5.7 and verified with the existing reflection smoke harness. The test keeps `UnrealAssets` as the default backend and verifies that the backend UI switch still toggles between curated assets and semantic proxies. For the descriptor path, the smoke harness loads a real SPPA-DESC-0.2 fixture from the atomic descriptor benchmark and calls the native actor API. The fixture `sppa-8d4e38132a285493` contains 12 parts; Unreal generated 12 static-mesh components and attached `SPPA_MATERIAL_ROLE`, `SPPA_EVIDENCE_SOURCE`, and `SPPA_UNCERTAINTY_STYLE` component tags. The same smoke verified that invalid JSON is rejected without changing the current proxy, that descriptor reconfiguration changes the part count, and that a pose-update packet is accepted for an existing descriptor proxy without rebuilding topology. A pose-update packet with a mismatched `descriptor_id` is rejected by the actor-level smoke test.

This closes the earlier implementation gap between the Python descriptor and the Unreal semantic-proxy backend. It does not establish packaged-build performance, VR frame rate, draw calls, GPU cost, dense-scene behavior, or operator benefit.

20.5 Unreal Actor-Level Microbenchmark

We then measured the native actor path in Unreal Editor-Cmd mode. The benchmark creates transient actors directly through the Unreal Python API, applies five pose updates per actor, applies one shape update, and destroys the actors. It uses the same descriptor fixtures saved by the Python descriptor benchmark and compares six baselines: no rendering, billboard label, box proxy, ellipsoid proxy, the legacy semantic proxy, and descriptor-driven SPPA. The run used 10, 50, 100, 250, and 500 actors with ten repetitions per count. The artifact directory is `experiments/sppa_unreal_backend/20260702T084924Z_editor_actor_microbenchmark`.

The manifest reports no exceptions and zero internal row failures for create, pose, shape, or no-op actions. All rows reused both component count and component names during pose and shape updates. The result is deliberately not presented as a latency victory over trivial boxes or ellipsoids: those baselines are faster and use far fewer components. The measured SPPA claim supported by

Table 15: Preliminary Unreal Editor-Cmd actor microbenchmark. Times are per-object p50/p95 milliseconds across 50 rows per method. The p95 values are descriptive rather than a final latency guarantee. This is not a packaged build, not a `/api/ui/data` replay, not render-thread timing, and not VR frame-time evidence.

Method	Components at 500	Create P50	Create P95	Pose P50	Pose P95	Shape P50	Shape P95
<code>no_render</code>	0	0.0029	0.0034	0.00009	0.00012	0.00000	0.00002
<code>billboard_label</code>	0	0.2915	0.4184	0.0087	0.0102	0.00012	0.00026
<code>box_proxy</code>	500	0.3619	0.4619	0.0120	0.0128	0.0091	0.0102
<code>ellipsoid_proxy</code>	500	0.3456	0.4518	0.0087	0.0094	0.0054	0.0065
<code>legacy_semantic_proxy</code>	2180	0.6110	0.8250	0.0149	0.0195	0.0135	0.0216
<code>sppa_desc</code>	2680	0.9689	3.9896	0.0287	0.0503	0.0316	0.0502

this microbenchmark is narrower: a richer semantic part actor can be created around 1 ms per object at the median in this Editor-Cmd setup and updated without component reconstruction at roughly 0.03 ms per object at the median. The SPPA create-time p95 is much higher than the median in this run, so packaged replay, warm/cold cache separation, and frame-time measurement are still required. Whether the additional semantic structure changes operator recognition or remains acceptable in a packaged dense VR scene is not shown by this microbenchmark.

20.6 Component Payload Replay

To move one step closer to the actual telemetry lifecycle, we added a debug entry point on the Unreal telemetry component that applies a JSON payload with the same `obstacles[]` shape used by `/api/ui/data`, but without HTTP. This benchmark exercises component parsing, entity upsert, backend selection, actor spawning, descriptor/update packet routing, pose updates, and shape updates. It compares the default asset backend against the semantic-proxy backend using the same obstacle payload. In this preliminary run the asset backend uses `StaticMeshActor` placeholders rather than the project curated asset library, so it is a lifecycle baseline, not a final visual-asset comparison.

The run used 10, 50, 100, 250, and 500 actors, ten repetitions per count, and five pose-update batches per actor. The run identifier is `20260702T085905Z`; artifacts are stored under the component-replay experiment directory. The benchmark order was not randomized, cold and warm cache behavior were not separated, and p95 values should be read as descriptive preliminary tails from ten repetitions per size rather than as operational tail latency guarantees.

Table 16: Preliminary component payload replay using the `/api/ui/data`-style obstacle payload shape. Times are per-object p50/p95 milliseconds across 50 rows per backend, with per-size summaries retained as artifacts. This bypasses HTTP and does not measure packaged-build, render-thread, GPU, draw-call, or VR frame-time cost.

Backend	Components at 500	Create P50	Create P95	Pose P50	Pose P95	Shape P50	Shape P95
<code>UnrealAssets</code> placeholder	500	0.5116	0.9882	0.0226	0.0447	0.0323	0.0703
<code>SemanticProxy/SPPA-DESC</code>	2680	1.2892	1.6423	0.0426	0.0531	0.0683	0.0829

The manifest reports no failures. This result supports only a narrow lifecycle claim: the same obstacle payload shape can drive both backends, and the SPPA component path remains low-millisecond per actor for creation and below 0.1 ms per actor for pose/shape updates at p95 in this Editor-Cmd replay. The table aggregates all tested scene sizes; per-size summaries and component-identity traces are stored as benchmark artifacts. The SPPA path also used many more components than the placeholder asset path (2680 vs. 500 at 500 actors), so the benchmark must not be interpreted as render scalability or visual-quality equivalence. It does not settle the final runtime question because it bypasses HTTP; the next subsection closes only the local polling-path gap, while the real curated assets, packaged execution, render-thread/GPU work, and VR presentation remain unmeasured.

20.7 HTTP Poll Replay Through the Unreal Component

We next exercised the Unreal HTTP polling branch with a local loopback server that served the same `/api/ui/data` obstacle payload shape. The component used a debug-only blocking helper, `PollNowBlockingForTest`, so the run could deterministically wait for each Unreal HTTP-module response inside `Editor-Cmd`. Normal runtime polling remains asynchronous. This benchmark covers HTTP request dispatch, response parsing, component batch ingestion, backend selection, actor spawning, descriptor/update packet routing, pose updates, and shape updates. It still does not cover a live flight server, real network jitter, packaged execution, render-thread/GPU work, draw calls, or VR presentation.

The valid run identifier is `20260702T101020Z`; artifacts are stored in the HTTP poll replay experiment directory. The run loaded 52 descriptor fixtures with input hash `459d477d66f5b9af`, used 10, 50, 100, 250, and 500 actors, ten repetitions per size, and five pose-update batches per actor. The manifest reports no failures; the row artifacts contain 1050 HTTP GET requests. The run uses seed `20260702` and interleaves backend order within each count/repetition group through a rotated/reversed counterbalanced schedule. The schedule and actual execution order are stored as `http_poll_replay_schedule.jsonl` and `http_poll_replay_order_trace.csv`. Benchmark garbage collection is requested before each condition, but cold and warm cache behavior are still not separated. In addition to the two operational spawn backends, the run includes a `no_render` debug baseline that parses the same HTTP payloads and updates component entity state while disabling actor spawning through a test-only flag. This baseline is not exposed as an operational UI backend. The run also emits paired incremental-delta artifacts comparing `UnrealAssets` and `SemanticProxy` against the same `no_render` component-lifecycle baseline and against each other. The same run emits `sppa_runtime_update_trace.csv`, `sppa_runtime_update_trace.jsonl`, and `sppa_runtime_update_summary.json`, which record resolver match type, scale/material source, uncertainty tags, descriptor identity, component reuse, and component creation/destruction counts around each create, pose-update, shape-update, or no-op payload. An earlier run that accidentally loaded only one descriptor fixture was discarded and is not used in the reported numbers.

Table 17: Preliminary local-loopback HTTP poll replay through the Unreal telemetry component. Times are per-object p50/p95 milliseconds across 50 rows per backend. This is an `Editor-Cmd` replay through the Unreal HTTP module, not a packaged-build, render-thread, GPU, draw-call, or VR frame-time benchmark.

Backend	Components at 500	Create P50	Create P95	Pose P50	Pose P95	Shape P50	Shape P95
NoRender debug	0	0.6577	1.3239	0.2209	0.6223	0.1425	0.5688
UnrealAssets placeholder	500	0.7129	1.3484	0.2681	0.6947	0.1477	0.6172
SemanticProxy/SPPA-DESC	2680	1.4897	2.4012	0.4068	0.8706	0.2250	0.6963

Table 18: Batch-level local-loopback HTTP poll replay by scene size. Create and shape columns report one HTTP poll with ten repetitions per size/backend. Pose columns report raw pose-update HTTP polls, five per repetition, for 50 pose polls per size/backend. Values are p50/p95 milliseconds; p95 is descriptive rather than a robust tail-latency estimate.

Backend	Actors	Create poll	Pose poll	Shape poll
NoRender	10	11.8/13.5	5.7/6.8	5.6/7.5
NoRender	50	32.9/36.2	12.3/14.3	8.9/10.4
NoRender	100	63.5/67.1	20.4/23.2	14.2/15.9
NoRender	250	159.2/163.0	45.5/48.6	29.9/32.0
NoRender	500	329.8/367.0	87.9/104.1	55.8/65.1
UnrealAssets	10	12.1/14.6	6.1/7.2	5.9/6.4
UnrealAssets	50	34.1/39.1	12.2/14.7	9.9/11.7
UnrealAssets	100	69.5/72.7	21.5/23.7	14.8/16.9
UnrealAssets	250	172.4/187.2	47.1/50.6	31.4/33.7
UnrealAssets	500	357.4/371.6	91.5/96.8	58.5/66.4
SemanticProxy	10	22.7/24.7	6.8/8.3	6.8/7.6
SemanticProxy	50	74.2/77.3	17.4/19.0	13.5/14.1
SemanticProxy	100	146.2/150.7	29.4/31.8	22.1/30.1
SemanticProxy	250	367.4/391.7	68.0/72.3	49.0/52.4
SemanticProxy	500	738.6/839.3	134.0/142.3	95.4/112.5

Table 19: Paired incremental batch deltas at 500 actors from the same local-loopback HTTP replay. Entries report p50 delta and mean delta with bootstrap 95% confidence interval in milliseconds. Negative or near-zero deltas in the **UnrealAssets** rows indicate measurement noise around the shared HTTP/JSON/entity-state path; the table should be read as preliminary component-path overhead, not as stage timing or rendering cost.

Comparison	Create Δ		Pose Δ		Shape Δ	Δ components
UnrealAssets-NoRender	24.5	[28.7; 9.3, 45.5]	1.1	[2.1; -1.0, 4.7]	0.9 [2.9; 0.4, 5.8]	500
SPPA-NoRender	413.9	[424.8; 401.1, 452.2]	45.5	[44.8; 42.2, 47.7]	40.1 [40.8; 37.0, 44.5]	2680
SPPA-UnrealAssets	387.8	[396.1; 371.7, 419.8]	40.8	[42.7; 40.3, 45.0]	37.9 [37.8; 35.2, 40.5]	2180

Table 20: Runtime update trace for the **SemanticProxy** backend in the same HTTP poll replay. Counts aggregate entity-level trace rows across all scene sizes and repetitions. The trace is actor/component instrumentation, not render-thread or GPU timing.

Phase	Rows	Created	Reused	Destroyed	Fallback	Yaw ambig.	Scale-source	rows
Create	2,950	24,740	0	0	550	2,950	690	metric / 2,260
Pose update	45,500	0	246,700	0	2,750	45,500	34,200	metric / 11,300
Shape update	9,100	0	49,340	0	550	9,100	6,840	metric / 2,260

This result is intentionally interpreted narrowly. It shows that the same HTTP-pollled obstacle payload can drive both backends and that descriptor-driven SPPA actors can be created and updated through the Unreal component without reported request or parsing failures in a local Editor-Cmd replay. It does not show a latency advantage over asset spawning, because the placeholder asset backend is faster for creation and uses fewer components. The no-render row also shows that large payload parsing and entity-state bookkeeping already consume a substantial part of the batch cost, especially at 500 actors. The SPPA-specific increment is therefore not the whole HTTP-poll cost, but the paired 500-actor deltas show an additional median create cost of 387.8 ms relative to the placeholder asset backend and an additional median pose-poll cost of 40.8 ms before rendering is measured. The new runtime trace supports a narrower but important implementation claim: pose and shape payloads update resident components without topology reconstruction, and resolver/fallback/yaw uncertainty metadata remain visible as actor tags. This table still does not show dense-scene render viability: at 500 actors, one SPPA create poll is hundreds of milliseconds in this replay, and one pose-update poll is over 100 ms before any render-thread or GPU cost is

measured. The 500-actor SPPA case also contains 2680 generated components. Packaged runtime, real curated assets, networked telemetry, render-thread/GPU timing, and operator interpretation are outside this replay.

20.8 Packaged Internal Render Benchmark

We then added an opt-in packaged benchmark runner to the Unreal plugin. The runner is activated only by the command-line flag `-SPPAPackagedBenchmark`; normal application maps and telemetry paths are otherwise unchanged. The benchmark packages the project, launches the executable, loads a dedicated empty benchmark map, and applies synthetic obstacle batches through the telemetry component batch API. This is an internal packaged replay with rendered frames, not live HTTP telemetry and not VR headset execution. The clean 2026-07-02 run is stored under run identifier `20260702T122743Z_packaged_render`; the packaging smoke run identifier is `20260702T122702Z_packaged_render`. Full artifact paths and commands are recorded in `docs/sppa_packaged_render_benchmark_20260702.md`.

Table 21: Packaged Unreal internal render benchmark, 100 synthetic obstacles, three repetitions, 30 warmup frames and 120 measured frames per phase. Frame times are Tick delta p50/p95 milliseconds. Draw calls and triangles are estimated from static mesh components, not GPU profiler counters. The `unreal_assets` backend uses a simple packaged placeholder actor, not the curated project asset library.

Backend	Components	Triangles	Draws	Create frame P95	Pose frame P95	Shape frame P95	Create apply P50
<code>no_render</code>	0	0	0	4.072	4.140	4.176	0.974
<code>unreal_assets</code>	100	4,800	100	4.414	5.227	4.810	5.952
<code>semantic_proxy</code>	432	208,864	432	4.817	17.814	19.256	46.067

The benchmark produced no benchmark-level failures, and the dedicated map avoided unrelated `BP_OpenSkyManager` HTTP polling observed in an earlier smoke attempt. Across the tested 10, 50, and 100 object densities, all SPPA phases stayed below 33.3 ms p95 in this packaged desktop run. However, the 100-object SPPA pose and shape update streams exceeded a 16.7 ms p95 frame budget, and SPPA had much higher create-apply cost than the placeholder asset backend because it builds many more primitive components. Therefore the allowed claim is limited: the current descriptor-driven semantic proxy can run in a packaged executable for 100 synthetic obstacles under a 30 FPS p95 budget on the tested desktop, but it is not yet a demonstrated VR/headset, live-network, or curated-asset result. Unreal CSV profiler support was also smoke-tested. The initial small profiler run produced a zero-byte CSV, but the later packaged directory contains two non-empty timestamped CSV profiler files. A separate parser extracted global counters from those files; the larger profile reported `FrameTime` p50/p95 4.662/6.318 ms, `GPTime` p50/p95 1.455/1.672 ms, `RHI/DrawCalls` p50 106, and local GPU memory p50 3,181 MB. These counters are not used for backend or phase claims because the CSV contains no SPPA backend/phase events beyond generic PSO events.

The instanced HISM backend was then optimized to defer render-state dirty marking during batch instance updates and flush once per touched group. The post-change packaged run `20260703T031246Z_packaged_render` used only the HISM backend at 100 objects, with three repetitions, 30 warmup frames, and 120 measured frames per phase. It produced no benchmark-level failures. The HISM route used 11 active draw groups; create/pose/shape frame P95 was 5.553/13.884/29.715 ms, and shape-frame P99 was 35.021 ms. Six of the 360 measured shape frames exceeded 33 ms. This initially suggested a narrow desktop 30 FPS P95 result for 100 synthetic objects, but subsequent no-CSV-profiler reruns did not reproduce the shape-update result. The dense sweep `20260703T033655Z_packaged_render` tested 100, 250, and 500 objects: shape-stream P95 was 38.657, 94.018, and 186.113 ms, respectively. An isolated 100-object rerun, `20260703T033840Z_packaged_render`, reported 37.126 ms shape-stream P95. The current supported conclusion is therefore that HISM batching substantially reduces draw-group pressure, not

that stable 30 FPS P95 shape-update performance has been demonstrated. It is still not a 60 FPS, headset, live-network, or phase-aligned GPU-profiler result.

The next regression added explicit partial updates to the HISM path. A payload may now set `partial_obstacle_update=true`; unmentioned entities are retained, while a mentioned entity first removes its previous instance keys and then upserts the new parts. The backend verifier covers this contract by creating two truck entities, updating only one with a reduced two-part `shape_param_update`, and checking that the live/HISM instance count is 6/6: two updated parts plus the four untouched parts of the second entity.

Three packaged runs then separate dense updates from changed-track scheduling. The rebuilt package 20260703T040154Z repeated the 100-object dense HISM case and still failed the 30 FPS P95 shape budget: create, pose, and shape frame P95 were 8.945, 21.315, and 46.151 ms, with 24 shape frames above 33 ms and four above 50 ms. A second run, 20260703T040406Z, kept pose updates dense but limited shape updates to 10% of objects. Shape-stream P95 was 9.559 ms at 100 objects, 15.477 ms at 250 objects, and 29.305 ms at 500 objects. Dense pose-stream P95 still failed at 250 and 500 objects: 52.474 and 104.502 ms. A third run, 20260703T040602Z, limited both pose and shape updates to 10%. Under that synthetic changed-track schedule, pose-stream P95 was 10.058 ms, 10.010 ms, and 16.528 ms for 100, 250, and 500 objects; shape-stream P95 was 9.629 ms, 15.454 ms, and 29.787 ms. No measured frame exceeded 33 ms in that run. This supports a partial-update scheduler contract only. It does not prove that 10% is a flight-derived update frequency, and it does not support dense all-object updates.

20.9 Preliminary Track-Lifecycle Measurement

To avoid overstating this policy, we measured only what the current artifacts can support. First, the current Python procedural template builder was executed 10,000 times per implemented class. Second, seven available controller `obstacle_ingest` logs were replayed offline to derive policy-implied create, update, shape-update, and topology-regeneration events. These logs contain 436,652 recorded track observations and 5,546 observed tracks. Three of the logs are explicitly marked `SIMULATION/AUTONOMOUS`; several older logs do not record workflow metadata. Some logs also mark obstacle samples as truncated. The result is therefore a preliminary lifecycle audit, not a native Unreal frame-time trace and not a final real-flight validation.

Table 22: Preliminary SPPA lifecycle costs from the current Python template builder and available controller obstacle-track logs. Update costs are descriptor/state updates only; Unreal actor transform and render-thread costs are not measured in this table.

Metric	n	P50	P95	P99
In-memory proxy creation	60,000	177.2 μ s	403.7 μ s	663.5 μ s
Debug OBJ/MTL export path	600	1.076 ms	1.522 ms	1.745 ms
Policy/state update per track observation	436,652	0.4 μ s	0.8 μ s	1.3 μ s

Under the track-persistent policy, the replay produced 5,546 create events, 410,766 pose and confidence updates, 20,340 shape and scale parameter updates, and zero topology-regeneration events. This zero should not be overinterpreted. It means that, in the recorded samples and under the declared policy, no track changed class or archetype in a way that required rebuilding the proxy topology. If the current non-parametric file-export path were used naively, the 20,340 shape and scale changes would become a full-rebuild upper bound. The publishable claim is therefore narrower: SPPA can in principle amortize object creation over track lifetime, but a final paper still needs native Unreal instrumentation for actor create, update, reconfigure, despawn, and frame-time impact.

21 Planned Evaluation Protocol

For full experimental validation, SPPA must be evaluated as an operational representation, not as a visual asset generator. The central comparison is against representation choices that a UAV digital twin would realistically use: no rendering, 3D bounding boxes, generic capsules or ellipsoids, billboard labels, curated low-poly assets when available, procedural actors, occupancy style volumes, and SPPA.

The minimum experimental matrix is:

Runtime generation.

Compare 3D box, capsule/ellipsoid, billboard label, curated low-poly asset loading, procedural proxy, and SPPA under cold/warm starts. Report P50/P95 latency, descriptor bytes, triangles, and memory.

Dense scene.

Compare boxes, capsules, billboards, low-poly assets, instanced primitive proxies, and detailed meshes in the target Unreal/VR profile at 10/50/100/500 entities. Report FPS, CPU/GPU frame time, and draw calls.

Proportional adaptation.

Compare fixed templates, bbox-only scaling, curated assets, and SPPA on synthetic variants plus real or simulated UAV-view detections. Report length/width/height error, BEV IoU, volume over/under-estimation, and yaw error.

Operator recognition.

Compare boxes, capsules/ellipsoids, billboard labels, curated low-poly assets, and SPPA in a randomized recognition task. Report class-family accuracy, response time, and confusion matrices.

Operational navigation.

Compare no rendering, boxes, occupancy volumes, and SPPA in degraded-video VR navigation. Report near misses, clearance, task time, NASA-TLX, and SAGAT-style probes.

Fallback and uncertainty.

Compare no rendering, generic boxes, uncertainty shells, and SPPA fallback on unknown or ambiguous detections. Report containment rate, false-confidence events, operator interpretation, and unknown-class failure rate.

Ablation.

Remove silhouette, footprint, track smoothing, semantic part graph, and uncertainty display from SPPA. Report BEV IoU, yaw error, jitter, and recognition accuracy.

Bandwidth.

Compare compressed video, telemetry boxes, SPPA descriptors, and full meshes in a network replay or controlled link emulator. Report bytes/s, packet rate, retransmission cost, and stale-object rate.

A minimal publishable benchmark should include both synthetic scale variants and real or simulated UAV-view detections. Synthetic variants test whether the same semantic graph adapts to different proportions, such as a short truck versus a long truck. UAV-view detections test the harder case: imperfect bounding boxes, partial silhouettes, top-down perspectives, and uncertain yaw. The evaluation must report failure cases, including occlusion, truncated detections, ambiguous front/back orientation, unknown classes, and misleading confidence displays.

22 Threats to Validity

Primitive proxies can misrepresent fine geometry, pose, articulation, and object-specific hazards. A proxy cow or tractor is not a true reconstruction of the observed object. The visual style may also create false confidence: a clean proxy can look more certain than the underlying detector, georeference, or scale estimate. This must be tested with uncertainty visualization rather than assumed.

Top-down UAV viewpoints, occlusion, motion blur, image compression, and partial detections can remove the evidence needed for part placement or signed yaw. The yaw model therefore exposes yaw_ambiguous, but the thresholds and display behavior are specified but not calibrated. Template bias is another risk: the seven current classes are too small to support broad claims about arbitrary semantic labels.

Language-guided semantic compilation introduces hallucination and certification risk. A language model may propose an implausible part graph for rare or ambiguous labels. SPPA must therefore review language-derived mappings against a finite archetype library, dimension priors, polygon budgets, and bounded fallback rules before runtime use. The current offline review-gate artifact implements a first cache-admission policy and rejects candidate recipes that force fallback, use unknown primitives/material roles, exceed hard limits, or request runtime LLM use. This reduces the specification gap, but it does not certify future language-model outputs or validate operator interpretation.

Finally, runtime performance is hardware- and engine-dependent. Triangle count alone does not predict headset frame rate, draw-call pressure, material cost, or tracking comfort. Dense-scene headset behavior must be measured directly.

23 Evidence Needed For Full System Validation

SPPA cannot be claimed as an evaluated UAV/VR runtime system until the following evidence exists:

End-to-end representation loop.

The manuscript lacks live detections, calibrated or real-flight georeferencing, a real network server, VR headset execution, and dense descriptor-driven replay from semantic-telemetry tracks. The required artifact is a packaged Unreal replay that creates SPPA descriptors and actors from recorded or simulated detections.

Shared baseline benchmark.

The packaged loop currently uses no-render and a placeholder asset backend. The required artifact is a common benchmark comparing boxes, capsules, billboards, curated assets, occupancy volumes, procedural alternatives, and SPPA on latency, geometry, frame rate, draw calls, and recognition.

Real shape adaptation.

Current evidence covers descriptor-level dimensions, synthetic oriented masks, calibrated synthetic mask-to-shape rows, and Unreal descriptor ingestion. The required artifact is a dataset with ground-truth or simulated footprints and BEV IoU, volume error, yaw error, and mask/silhouette ablations.

VR runtime.

OBJ timing and packaged desktop Tick timing are not headset frame timing. The required artifact is a packaged Unreal/VR replay reporting backend, count, poll rate, action mix, FPS, GameThread, RenderThread, GPU time, draw calls, triangles, memory, hitches, actor count, component count, seed, and build hash.

Operator evidence.

Readability and situation-awareness claims require a controlled user study with recognition, response time, navigation outcome, NASA-TLX, and SAGAT-style probes.

Fallback evidence.

Generic uncertainty fallback may create false confidence. The required artifact is an unknown-class and ambiguous-detection study with containment, operator interpretation, and failure analysis.

Bandwidth comparison.

Descriptor and update-packet bytes are measured and a preliminary link-budget model compares SPPA with modeled boxes, video, and mesh transfer. The required artifact is still a network replay or link-emulation report measuring bytes/s, packet rate, staleness, and failure modes under the intended communication architecture.

24 Reproducibility Notes

The benchmark directories store command logs, event JSONL files, object-level CSVs, GPU snapshots, process snapshots, generated meshes, and rendered contact sheets. The descriptor/update contract benchmark directory stores CSV rows, JSONL descriptor samples, a JSON summary, and a Markdown report. The preliminary bandwidth/link-budget directory stores SPPA JSON packet byte counts and modeled box, mesh, and video payload comparisons; it is not a measured radio-link replay. The text/image-to-3D dependency worktrees and model caches are machine-local and git-ignored because they contain large repositories and model weights. A submission-ready artifact still needs a clean committed source state or archived patch hash, model hashes, environment lock files, and permanent copies of the exact source snapshots used for each generator. SF3D and SPAR3D remain unresolved in this local pass because one timed out during warm loading and the other required model access that was not available. These unresolved baselines are recorded as setup/access outcomes, not as quality failures.

The 6 GB GPU budget is implemented through PyTorch’s allocator fraction. It is a portable-flight stress profile, not a hard GPU partition and not a guarantee that Unreal, video decoding, tracking, and generation can all co-reside at that budget without native profiling. The benchmark also uses clean synthetic/proxy RGBA inputs for image-to-3D methods; degraded UAV crops, compression, occlusion, small object size, and motion blur are not covered by this benchmark.

25 Limitations

The current prototype is class-template based and does not infer volumetric scale from a real detection pipeline. It does not consume live YOLO tracks, instance masks, calibrated camera geometry, georeferenced footprints, or uncertainty estimates as an end-to-end flight loop. Unreal can now ingest an optional descriptor and construct primitive components from `parts[]`, and this path has been exercised through reflection/smoke tests, Editor-Cmd actor and component replays, local-loopback HTTP replay, a local georeferenced pose regression, and a packaged internal render benchmark. It has not been benchmarked in a VR headset, live-network telemetry loop, or curated-asset operational scene. The current timings measure debug OBJ export, Python descriptor construction, Unreal actor operations, and packaged synthetic replay; they do not yet establish final instanced primitive or headset runtime performance.

The paper also does not include a pilot or user study. Claims about operator readability, situation awareness, workload, or navigation usefulness remain hypotheses until evaluated. The specified fallback behavior is implemented as a Python/Unreal smoke-tested display fallback, but it is not yet shown to reduce operator error, false confidence, or operational risk.

The next step is to connect the proxy generator to YOLO detections, instance masks or bounding boxes, UAV pose, camera calibration, georeferenced placement, and dense Unreal runtime

benchmarking. The contribution becomes publishable only if it is evaluated under the intended operational constraints: remote piloting, degraded video links, semantic telemetry, and VR real-time rendering.

26 Conclusion

SPPA, as presented here, is a specification for a semantic proxy rendering layer for telemetry-driven UAV digital twins. Static asset libraries may be incomplete or operationally expensive for long-tail detections, and neural image-to-3D systems optimize asset fidelity rather than bounded low-polygon display. Primitive proxy assembly is a plausible operating point for this gap, but the current evidence is limited to a seven-class fixed-template OBJ debug prototype, a Python SPPA-DESC/SPPA-UPD v0.2 descriptor/update benchmark, and an optional Unreal descriptor-ingestion path with actor/component replays and a packaged synthetic render benchmark that preserves the asset backend as the default. It is not yet a live, dense UAV/VR runtime benchmark.

The manuscript therefore should not be read as an evaluated UAV/VR system. It defines a runtime contract, descriptor shape, bounded compiler concept, mathematical fitting target, and planned evaluation protocol. Publication as a full experimental contribution requires the implementation, benchmark, fallback, bandwidth, and operator-study artifacts listed above.

References

- [1] Iro Armeni, Zhi-Yang He, JunYoung Gwak, Amir R. Zamir, Martin Fischer, Jitendra Malik, and Silvio Savarese. 3d scene graph: A structure for unified semantics, 3d space, and camera. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 5664–5673, 2019.
- [2] Irving Biederman. Recognition-by-components: A theory of human image understanding. *Psychological Review*, 94(2):115–147, 1987.
- [3] Garrick Brazil, Abhinav Kumar, Julian Straub, Nikhila Ravi, Justin Johnson, and Georgia Gkioxari. Omni3d: A large benchmark and model for 3d object detection in the wild. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- [4] James H. Clark. Hierarchical geometric models for visible surface algorithms. *Communications of the ACM*, 19(10):547–554, 1976.
- [5] Xavier Decoret, Fredo Durand, Francois X. Sillion, and Julie Dorsey. Billboard clouds for extreme model simplification. *ACM Transactions on Graphics*, 22(3):689–696, 2003.
- [6] Mica R. Endsley. Situation awareness global assessment technique (sagat). In *Proceedings of the IEEE 1988 National Aerospace and Electronics Conference*, 1988.
- [7] Mica R. Endsley. Toward a theory of situation awareness in dynamic systems. *Human Factors*, 37(1):32–64, 1995.
- [8] Epic Games. Nanite virtualized geometry in unreal engine, 2024.
- [9] Epic Games. Hardware and Software Specifications for Unreal Engine. Unreal Engine documentation, 2026. Accessed 2026-07-01.
- [10] Elisabetta Fedele, Boyang Sun, Leonidas Guibas, Marc Pollefeys, and Francis Engelmann. Superdec: 3d scene decomposition with superquadric primitives. *arXiv preprint arXiv:2504.00992*, 2025.

- [11] Sandra G. Hart and Lowell E. Staveland. Development of nasa-tlx (task load index): Results of empirical and theoretical research. In Peter A. Hancock and Najmedin Meshkati, editors, *Human Mental Workload*, pages 139–183. North-Holland, Amsterdam, 1988.
- [12] Armin Hornung, Kai M. Wurm, Maren Bennewitz, Cyrill Stachniss, and Wolfram Burgard. Octomap: An efficient probabilistic 3d mapping framework based on octrees. *Autonomous Robots*, 34(3):189–206, 2013.
- [13] Shaofei Hu, Gengshan Yang, Shunsuke Saito Lu, et al. Sam 3d animal: Pan-category 3d animal reconstruction from a single image with shape prior. *arXiv preprint arXiv:2605.07604*, 2026.
- [14] R. Kenny Jones, Theresa Barton, Xianghao Xu, Kai Wang, Ellen Jiang, Paul Guerrero, Niloy J. Mitra, and Daniel Ritchie. Shapeassembly: Learning to generate programs for 3d shape structure synthesis. *ACM Transactions on Graphics*, 39(6), 2020.
- [15] Heewoo Jun and Alex Nichol. Shap-e: Generating conditional 3d implicit functions. *arXiv preprint arXiv:2305.02463*, 2023.
- [16] Zizhang Li, Dor Litvak, Ruining Li, Yunzhi Zhang, Tomas Jakab, Christian Rupprecht, Shangzhe Wu, Andrea Vedaldi, and Jiajun Wu. Learning the 3d fauna of the web. *arXiv preprint arXiv:2401.02400*, 2024.
- [17] Chen-Hsuan Lin, Jun Gao, Luming Tang, Towaki Takikawa, Xiaohui Zeng, Xun Huang, Karsten Kreis, Sanja Fidler, Ming-Yu Liu, and Tsung-Yi Lin. Magic3d: High-resolution text-to-3d content creation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- [18] David Marr and H. Keith Nishihara. Representation and recognition of the spatial organization of three-dimensional shapes. *Proceedings of the Royal Society of London. Series B, Biological Sciences*, 200(1140):269–294, 1978.
- [19] Pascal Mueller, Peter Wonka, Simon Haegler, Andreas Ulmer, and Luc Van Gool. Procedural modeling of buildings. *ACM Transactions on Graphics*, 25(3):614–623, 2006.
- [20] Alex Nichol, Heewoo Jun, Prafulla Dhariwal, Pamela Mishkin, and Mark Chen. Point-e: A system for generating 3d point clouds from complex prompts. *arXiv preprint arXiv:2212.08751*, 2022.
- [21] NVIDIA. GeForce RTX 50 Series Gaming Laptops: Compare Specs. Web page, 2026. Accessed 2026-07-01.
- [22] Helen Oleynikova, Zachary Taylor, Marius Fehr, Juan Nieto, and Roland Siegwart. Voxblox: Incremental 3d euclidean signed distance fields for on-board mav planning. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 1366–1373, 2017.
- [23] Open Geospatial Consortium. Ogc 3d tiles 1.1 community standard, 2023.
- [24] Despoina Paschalidou, Ali Osman Ulusoy, and Andreas Geiger. Superquadrics revisited: Learning 3d shape parsing beyond cuboids. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- [25] Ben Poole, Ajay Jain, Jonathan T. Barron, and Ben Mildenhall. Dreamfusion: Text-to-3d using 2d diffusion. *arXiv preprint arXiv:2209.14988*, 2022.
- [26] PyTorch Contributors. torch.cuda.memory.set_per_process_memory_fraction. PyTorch documentation, 2026. Accessed 2026-07-01.

- [27] Antoni Rosinol, Arjun Gupta, Marcus Abate, Jingnan Shi, and Luca Carlone. 3d dynamic scene graphs: Actionable spatial perception with places, objects, and humans. In *Robotics: Science and Systems*, 2020.
- [28] Franc Solina and Ruzena Bajcsy. Recovery of parametric models from range images: The case for superquadrics with global deformations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 12(2):131–147, 1990.
- [29] Stability AI. Introducing stable fast 3d: Rapid 3d asset generation from single images, 2024.
- [30] Stability AI. Sf3d: Stable fast 3d mesh reconstruction with uv-unwrapping and illumination disentanglement. *arXiv preprint arXiv:2408.00653*, 2024.
- [31] Jiaxiang Tang, Zhaoxi Chen, Xiaokang Chen, Tengfei Wang, Gang Zeng, and Ziwei Liu. Lgm: Large multi-view gaussian model for high-resolution 3d content creation. *arXiv preprint arXiv:2402.05054*, 2024.
- [32] Team Hunyuan3D. Hunyuan3d 2.1: From images to high-fidelity 3d assets with production-ready pbr material. *arXiv preprint arXiv:2506.15442*, 2025.
- [33] Team Hunyuan3D. Hy3d-bench: Generation of 3d assets. *arXiv preprint arXiv:2602.03907*, 2026.
- [34] Dmitry Tochilkin, David Pankratz, Zexiang Liu, Zixuan Huang, Adam Letts, Yangguang Li, Ding Liang, Christian Laforte, Varun Jampani, and Yan-Pei Cao. Triposr: Fast 3d object reconstruction from a single image. *arXiv preprint arXiv:2403.02151*, 2024.
- [35] Shubham Tulsiani, Hao Su, Leonidas J. Guibas, Alexei A. Efros, and Jitendra Malik. Learning shape abstractions by assembling volumetric primitives. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2635–2643, 2017.
- [36] Valve Corporation. Steam Hardware & Software Survey: May 2026. Web survey, 2026. Accessed 2026-07-01.
- [37] Zhengyi Wang, Yikai Wang, Yifei Chen, Chendong Xiang, Shuo Chen, Dajiang Yu, Chongxuan Li, Hang Su, and Jun Zhu. Crm: Single image to 3d textured mesh with convolutional reconstruction model. *arXiv preprint arXiv:2403.05034*, 2024.
- [38] Kailu Wu, Fangfu Liu, Zhihan Cai, Runjie Yan, Hanyang Wang, Yating Hu, Yueqi Duan, and Kaisheng Ma. Unique3d: High-quality and efficient 3d mesh generation from a single image. *arXiv preprint arXiv:2405.20343*, 2024.
- [39] Shun-Cheng Wu, Johanna Wald, Keisuke Tateno, Nassir Navab, and Federico Tombari. Scene-graphfusion: Incremental 3d scene graph prediction from rgb-d sequences. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021.
- [40] Jianfeng Xiang, Xiaoxue Chen, Sicheng Xu, Ruicheng Wang, Zelong Lv, Yu Deng, Hongyuan Zhu, Yue Dong, Hao Zhao, Nicholas Jing Yuan, and Jiaolong Yang. Native and compact structured latents for 3d generation. *arXiv preprint arXiv:2512.14692*, 2025.
- [41] Jianfeng Xiang, Zelong Lv, Sicheng Xu, Yu Deng, Ruicheng Wang, Bowen Zhang, Dong Chen, Xin Tong, and Jiaolong Yang. Structured 3d latents for scalable and versatile 3d generation. *arXiv preprint arXiv:2412.01506*, 2024.
- [42] Jiale Xu, Weihao Cheng, Yiming Gao, Xintao Wang, Shenghua Gao, and Ying Shan. Instantmesh: Efficient 3d mesh generation from a single image with sparse-view large reconstruction models. *arXiv preprint arXiv:2404.07191*, 2024.

- [43] Yikang Yang, Zhanpeng Hu, Youtian Lin, Mengqi Zhou, Jingxi Xu, Feihu Zhang, Jiaheng Liu, and Yao Yao. P3d-bench: Benchmarking mllms for parametric 3d generation and structural reasoning. *arXiv preprint arXiv:2606.11152*, 2026.
- [44] Jin Yao, Hao Gu, Xuweiyi Chen, Jiayun Wang, and Zezhou Cheng. Open vocabulary monocular 3d object detection. *arXiv preprint arXiv:2411.16833*, 2024.
- [45] Hang Yu, Yufei Xu, Jing Zhang, Wei Zhao, Ziyu Guan, and Dacheng Tao. Ap-10k: A benchmark for animal pose estimation in the wild. In *Advances in Neural Information Processing Systems*, 2021.
- [46] Zibo Zhao, Zeqiang Lai, Qingxiang Lin, Yunfei Zhao, Haolin Liu, et al. Hunyuan3d 2.0: Scaling diffusion models for high resolution textured 3d assets generation. *arXiv preprint arXiv:2501.12202*, 2025.